Microsoft

Operational best practices for Azure Kubernetes Service

(邦題: Azure Kubernetes Service (AKS) 管理のベスト プラクティス)

Saurya Das

Senior Program Manager
Azure Kubernetes Service
Microsoft Corporation.

CI32

Agenda

- Cluster Isolation and Resource Management
- Networking
- Securing your Environment
- Scaling your Applications and Cluster
- Logging and Monitoring

Cluster Isolation and Resource Management

Cluster Isolation Patterns: Physical Isolation

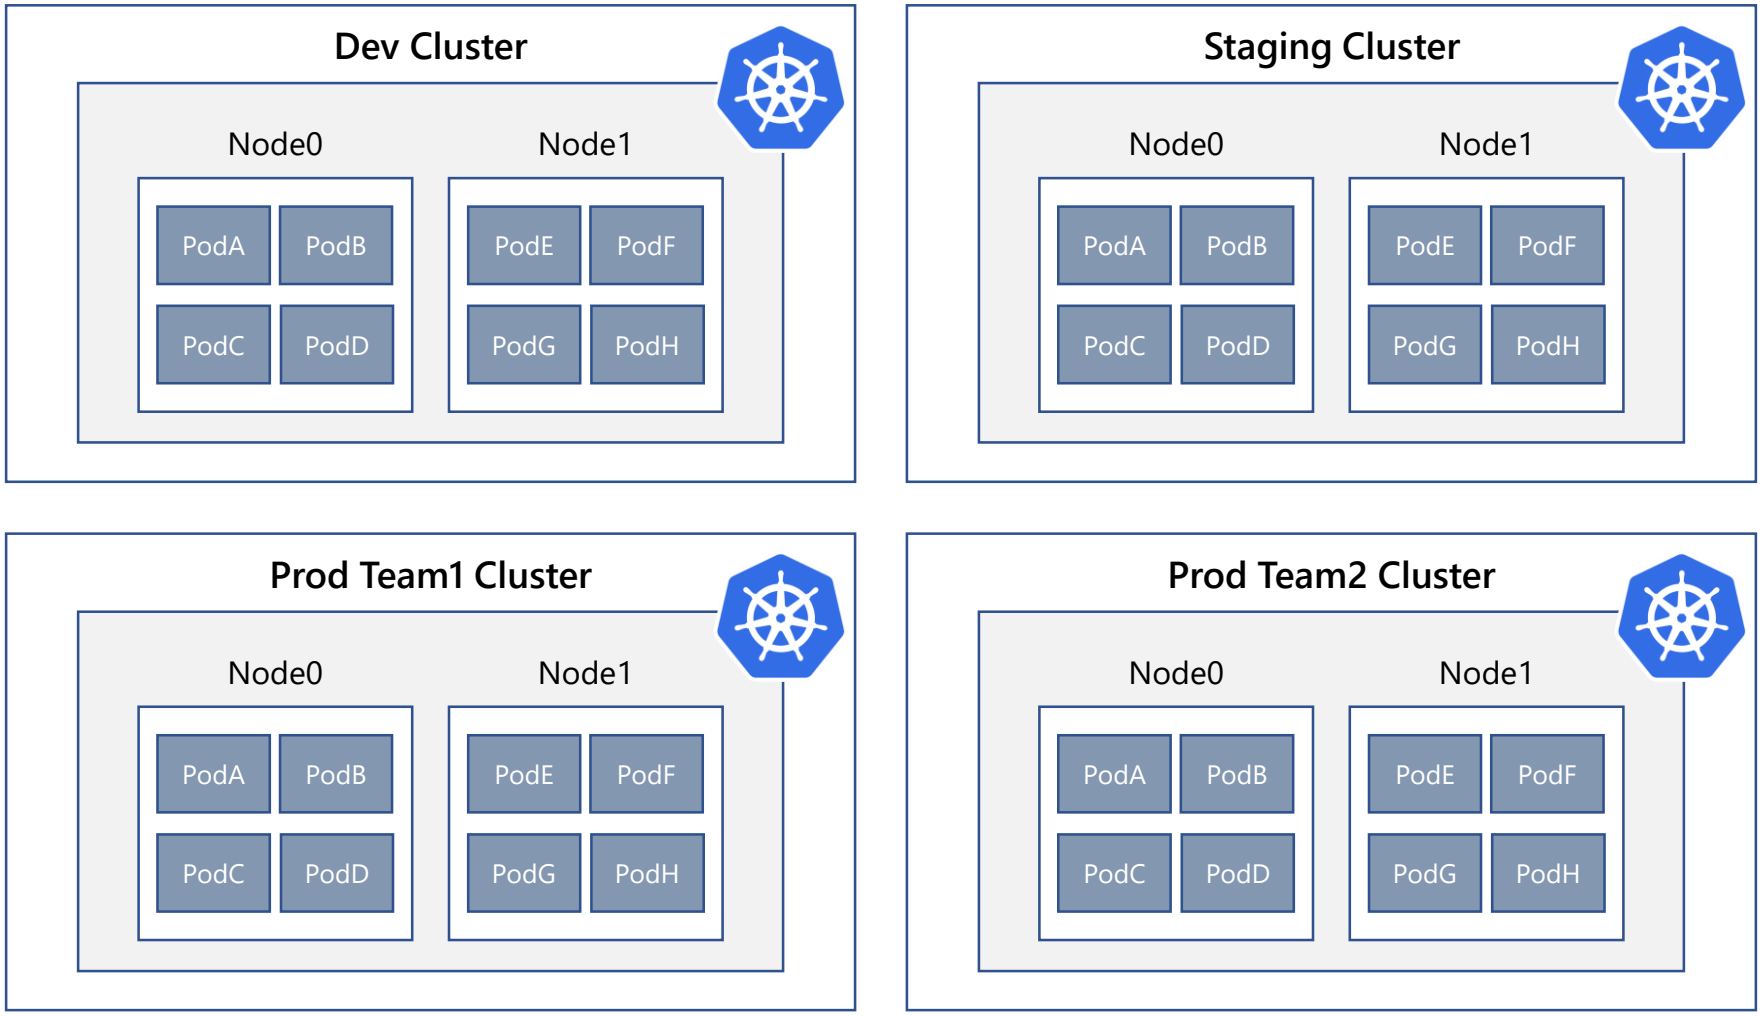

Cluster Isolation Patterns: Logical Isolation

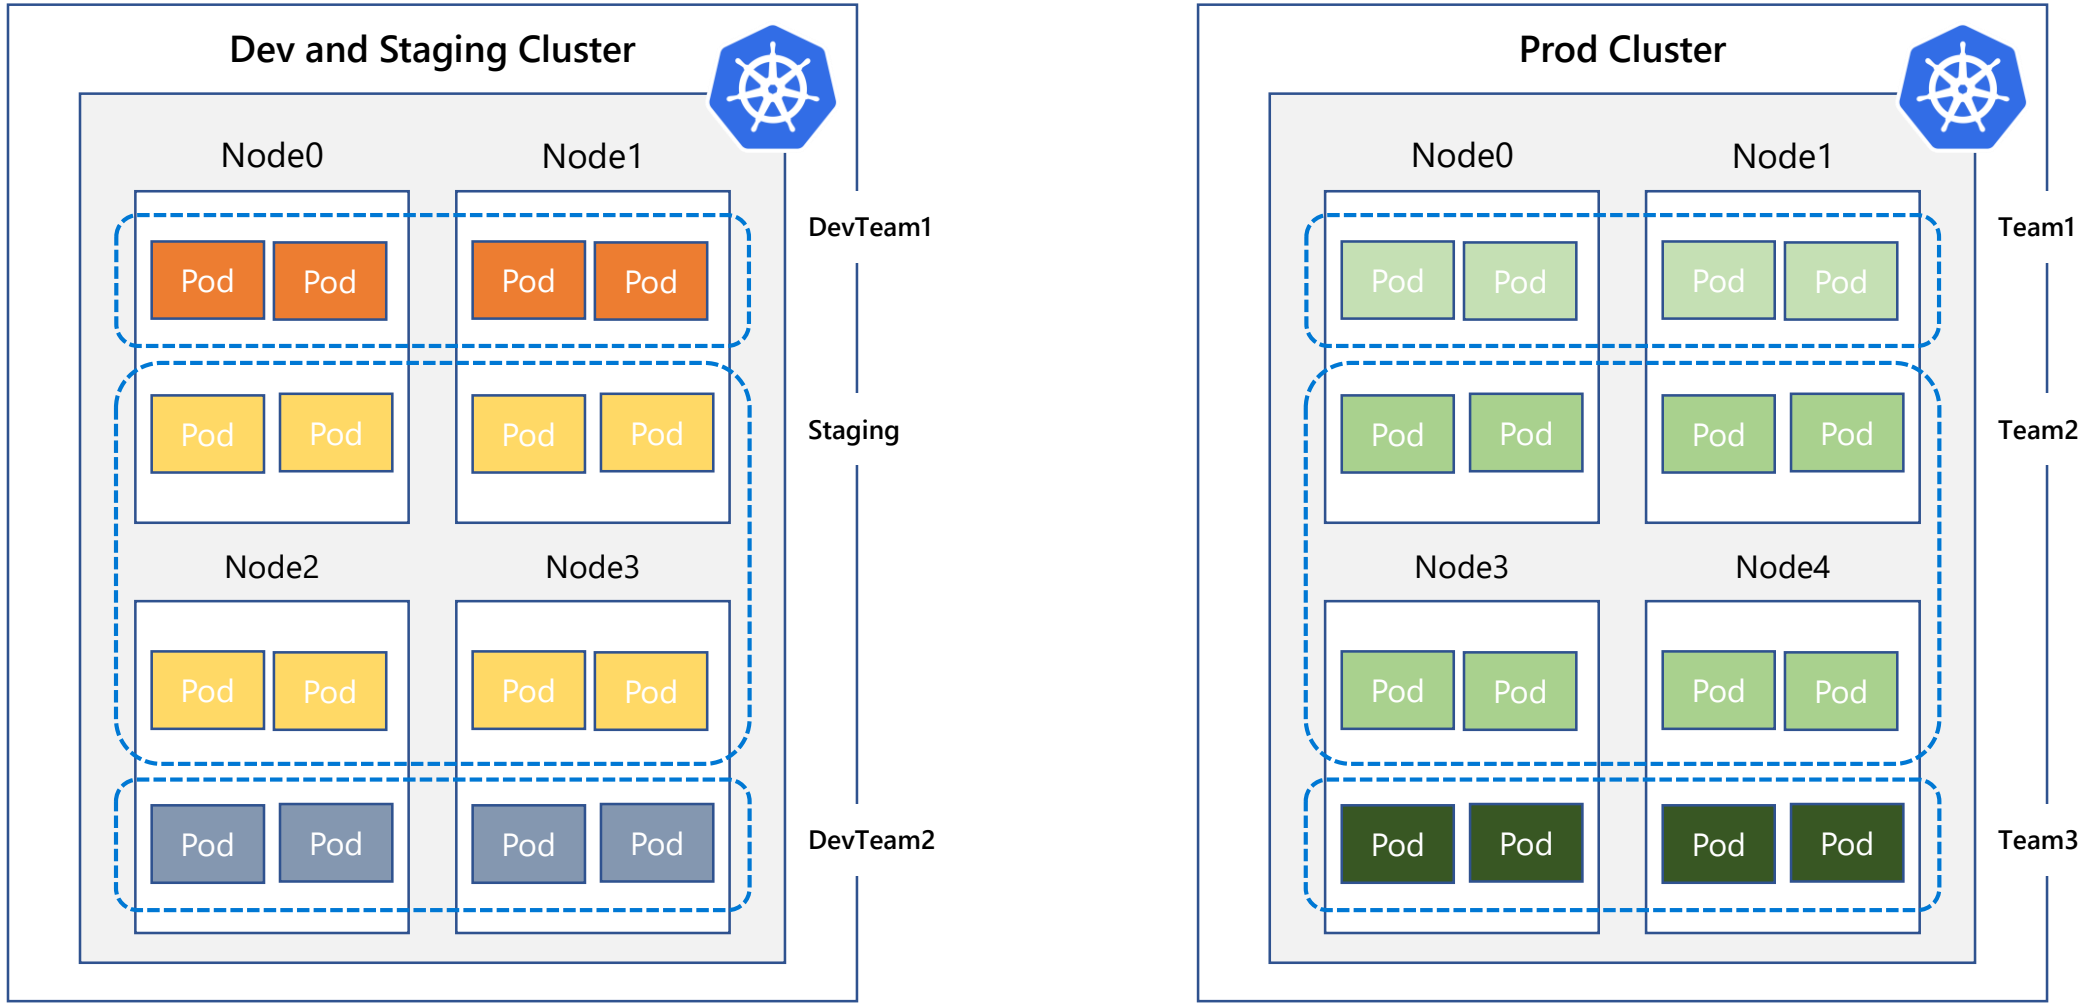

Kubernetes Namespaces

- **Namespaces** Object is the logical Isolation boundary
- Kubernetes has features to help us safely isolate tenants
  - **Scheduling**: Resource Quota
  - **Network** Isolation using Network Policies
  - **Authentication and Authorization**: RBAC and Pod Security Policy

- Note: **Container Level isolation** still need to be done to achieve hard Isolation

Kubernetes Resource Quotas

- Constraints that limit aggregate resource consumption per namespace
- You can limit Compute Resources (CPU,Memory, Strage,…) and/or limit the number of Objects (Pods, Services, etc..) and
- When enabled, users must specify requests or limits, otherwise the quota system will fail the request.
- Kubernetes will not overcommit

Create a namespace:

```
$ kubectl create namespace ignite
```

Apply a resource quota to the namespace:

admin/resource/ignite.yaml

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: mem-cpu-demo
spec:
hard:
    requests.cpu: "1"
    requests.memory: 1Gi
    limits.cpu: "2"
    limits.memory: 2Gi
```

Physical vs. Logical Isolation

	Physical	Logical
Pod Density	Low to Medium	Medium to High
Cost	$$	$
Kubernetes Experience	Low to Medium	Medium to High
Security	High (Surface is small)	High*
Blast Radius of Changes	Small	Big
Management and Operations	Owner Team	Single or Cross Functional Team

*Logical Isolation via Namespaces can achieve hard isolation assuming the cluster admin has applied all the required security controls

Demo

Create namespace
Apply a resource quota
Deploy a simple container within limits
Deploy another container beyond limits that fails

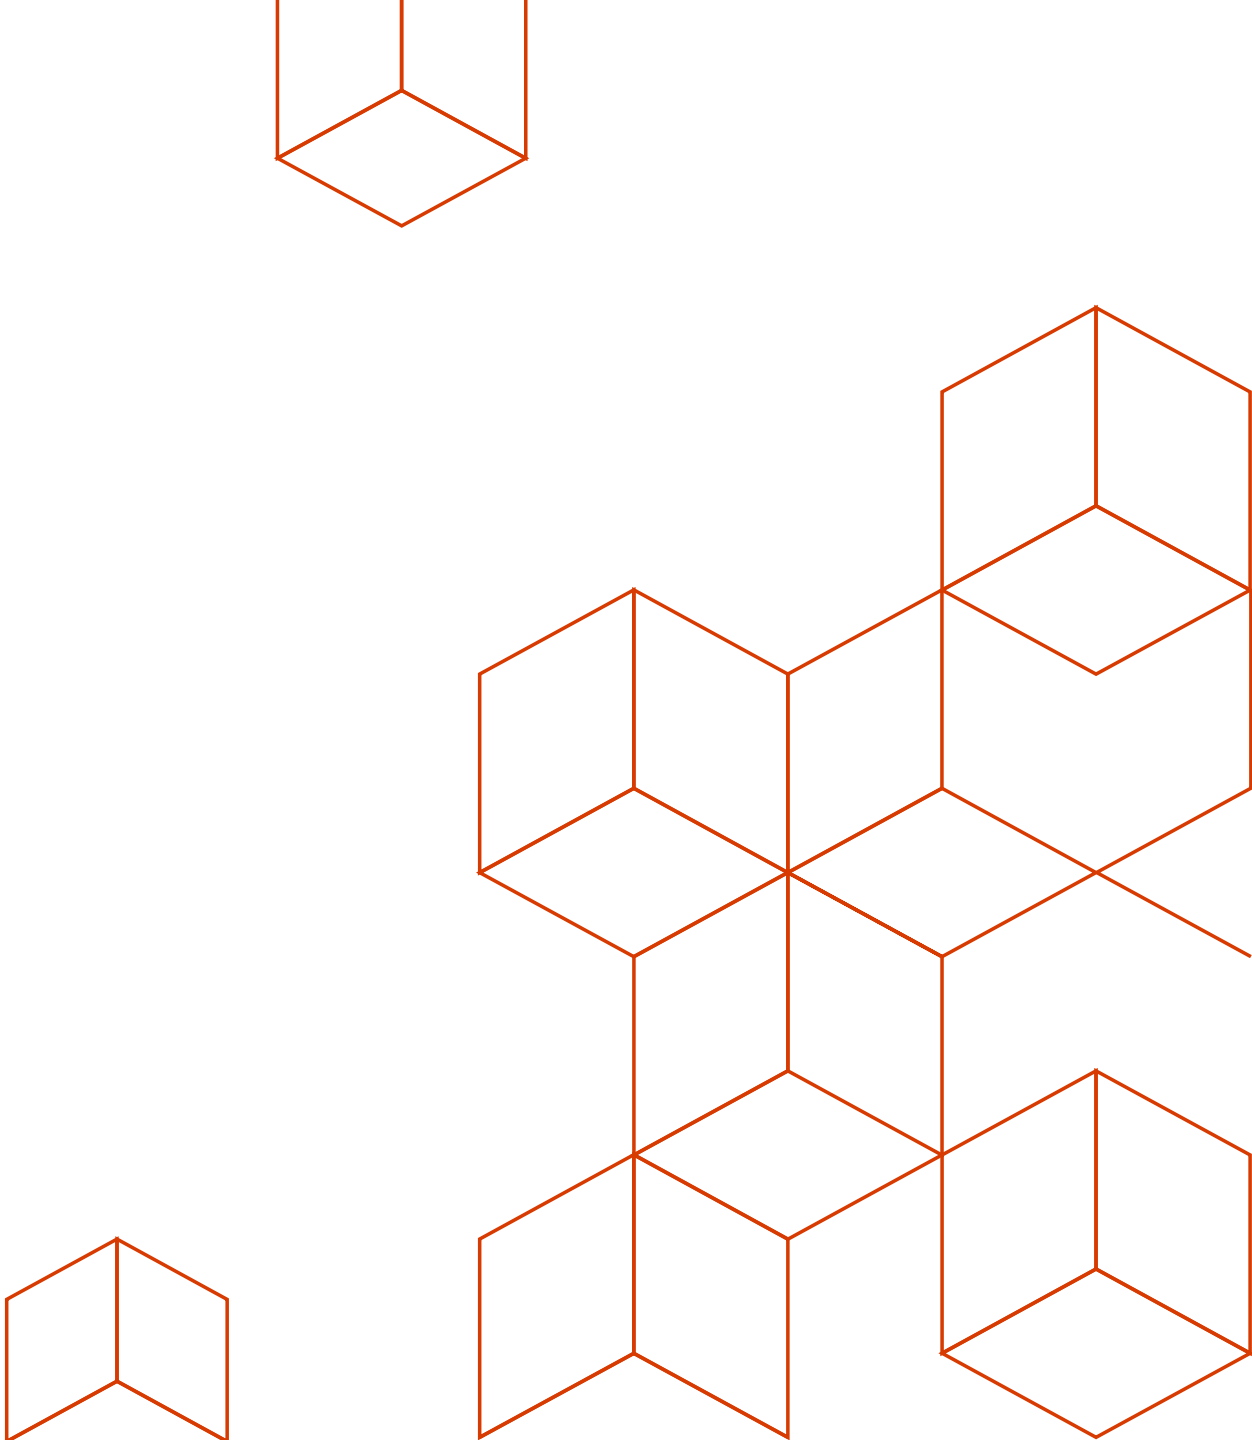

Kube-advisor

- Diagnostic tool for Kubernetes clusters. At the moment, it returns pods that are missing resource and request limits.
- More info can be found at https://github.com/Azure/kube-advisor

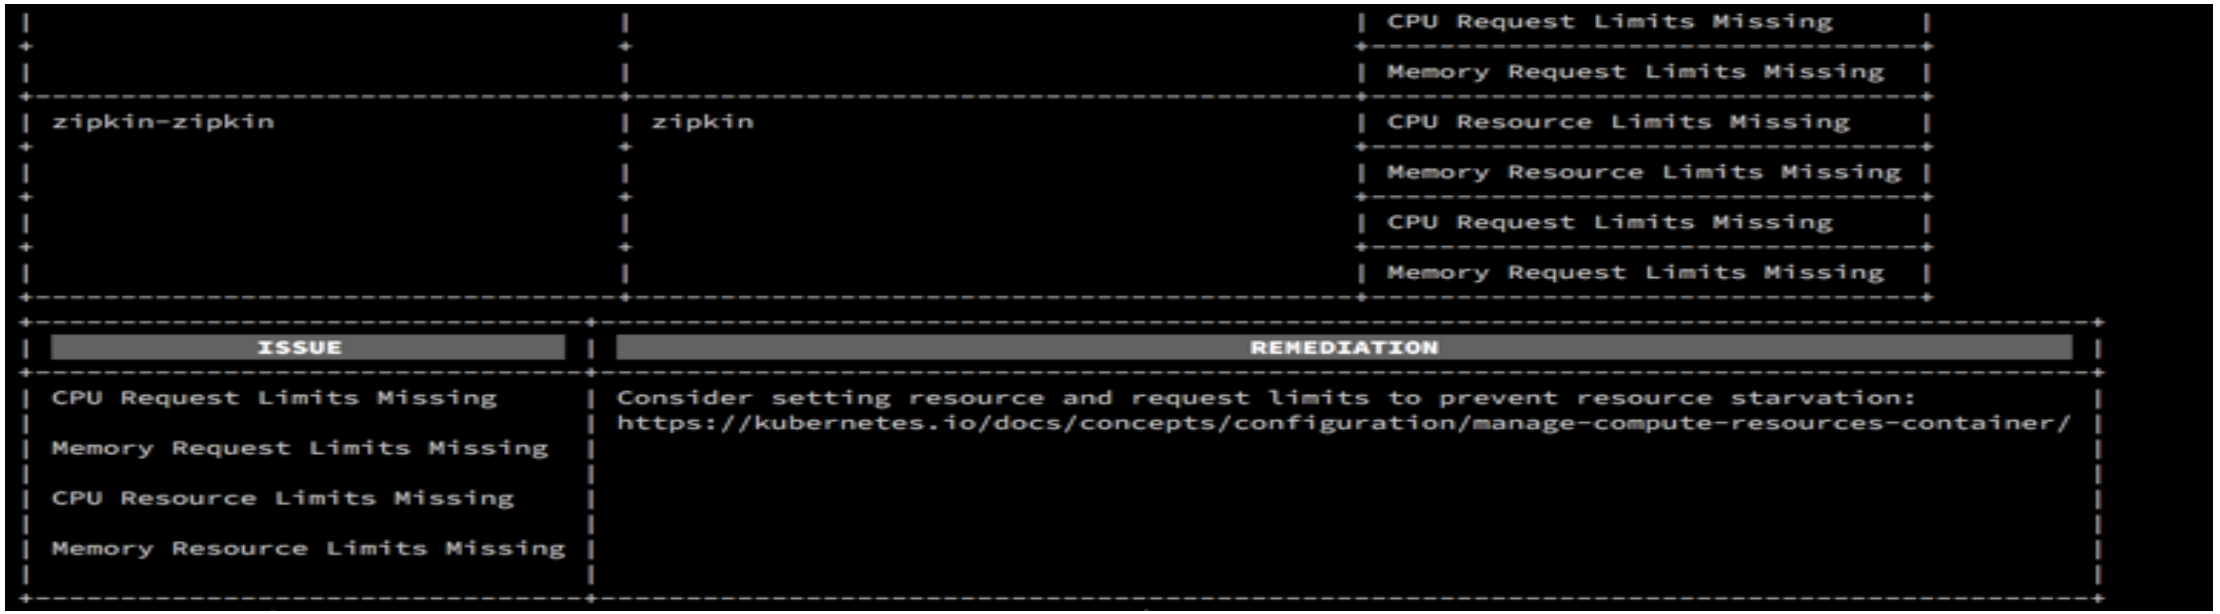

		CPU Request Limits Missing
		Memory Request Limits Missing
zipkin-zipkin	zipkin	CPU Resource Limits Missing
		Memory Resource Limits Missing
		CPU Request Limits Missing
		Memory Request Limits Missing

ISSUE	REMEDIATION
CPU Request Limits Missing Memory Request Limits Missing CPU Resource Limits Missing Memory Resource Limits Missing	Consider setting resource and request limits to prevent resource starvation: https://kubernetes.io/docs/concepts/configuration/manage-compute-resources-container/

VS Code extension for warnings

- Kubernetes VS Code extension adding warnings for resource request/limits

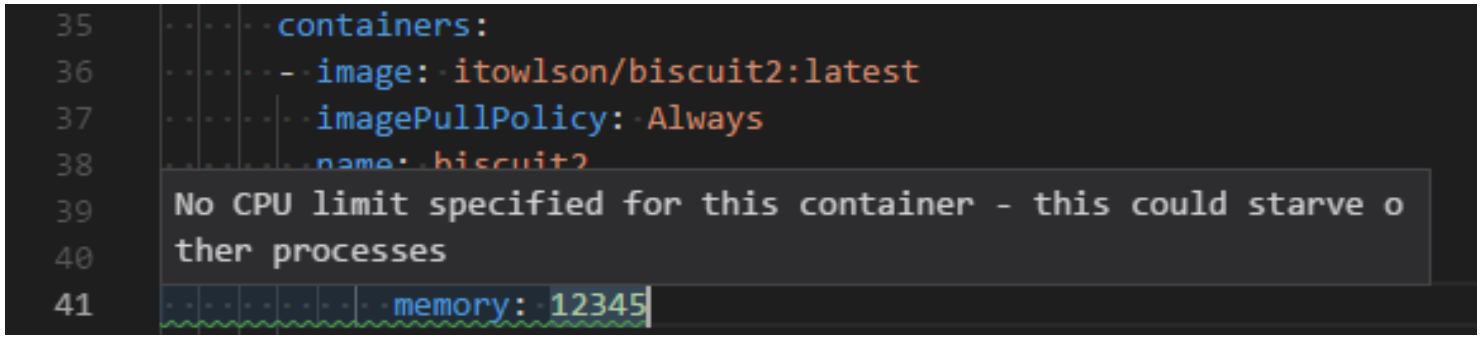

Cluster Isolation - Summary

- Think of the sensitivity of the workload, cost, organization culture, operations model, and blast radius, when trying to choose which isolation pattern to use, a mixture is fine too.

- Always use Namespaces even in physical isolation, never use the Default Namespace for production workloads

- Apply Resource Quotas

Networking

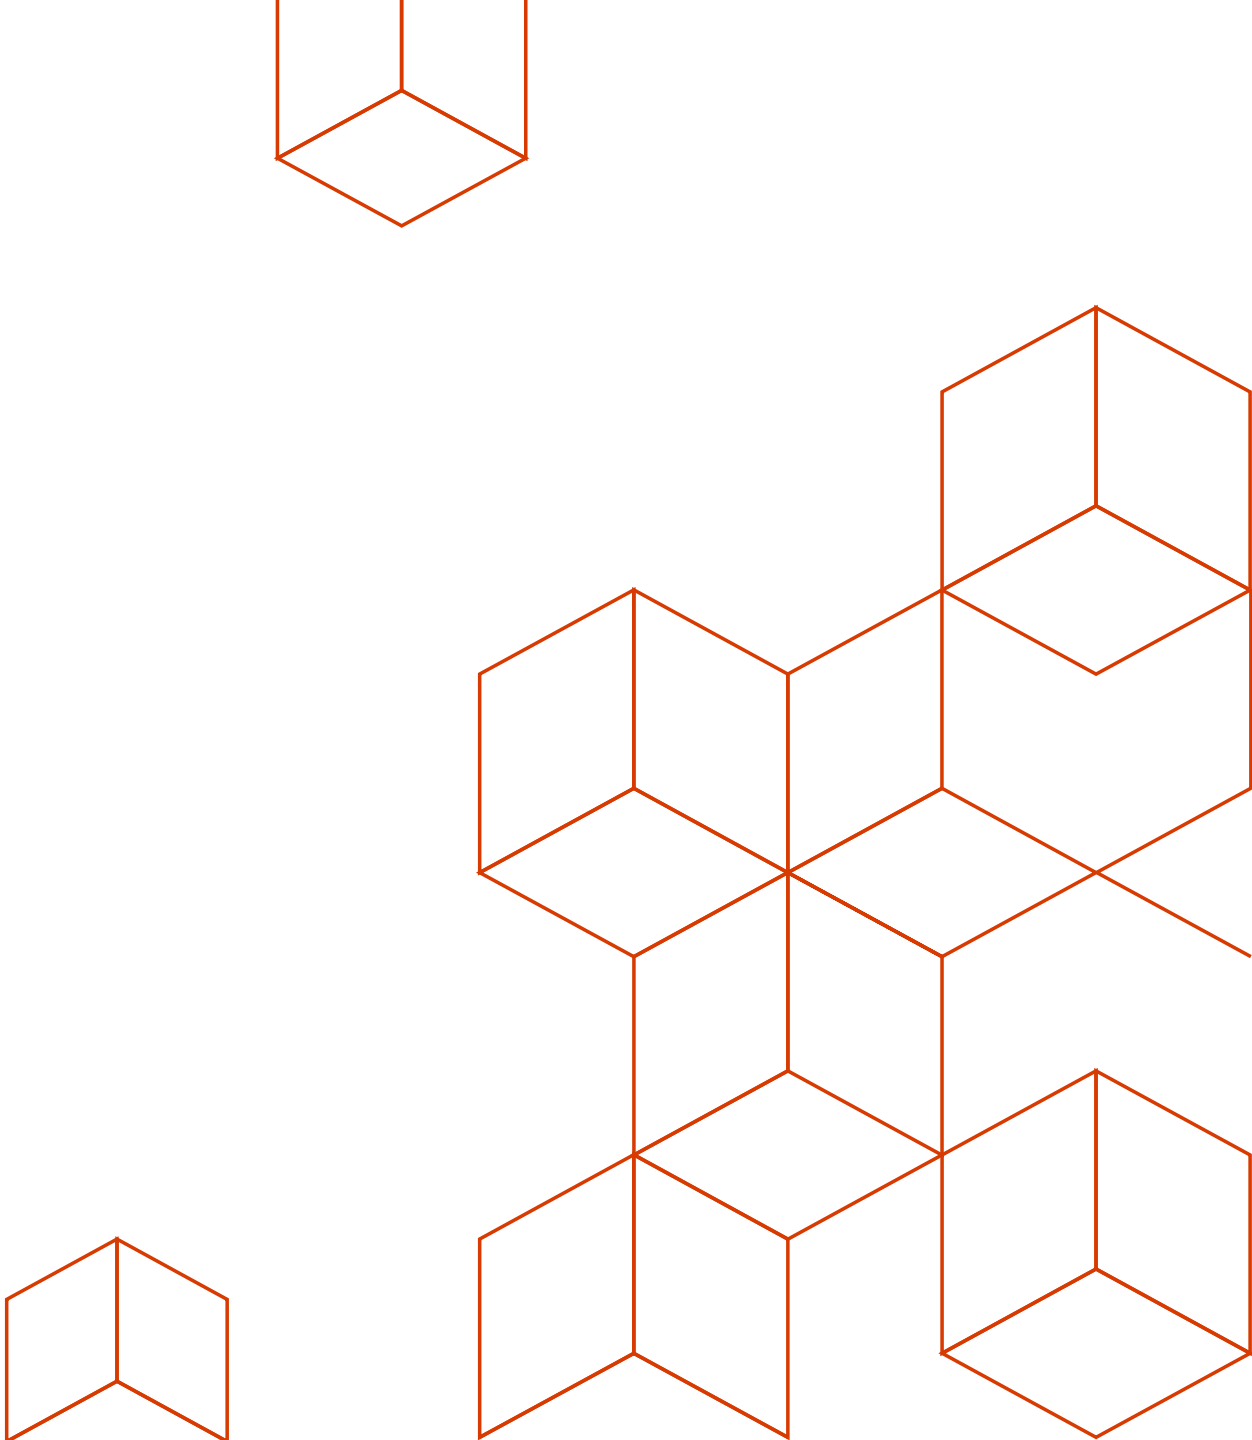

AKS Basic Networking

- Done using **Kubenet** network plugin and has the following features
  - Nodes and Pods are placed on **different** IP subnets
  - User Defined Routing and IP Forwarding is for connectivity between Pods across Nodes.

- Drawbacks
  - 2 different IP CIDRs to manage
  - Performance impact
  - Peering or On-Premise connectivity is hard to achieve

AKS Advanced Networking

- Done using the Azure CNI (Container Networking Interface)
  - **CNI** is a vendor-neutral protocol, used by container runtimes to make requests to Networking Providers
  - **Azure CNI** is an implementation which allows you to integrate Kubernetes with your VNET

- Advantages
  - Single IP CIDR to manage
  - Better Performance
  - Peering and On-Premise connectivity is out of the box

AKS with Advanced Networking

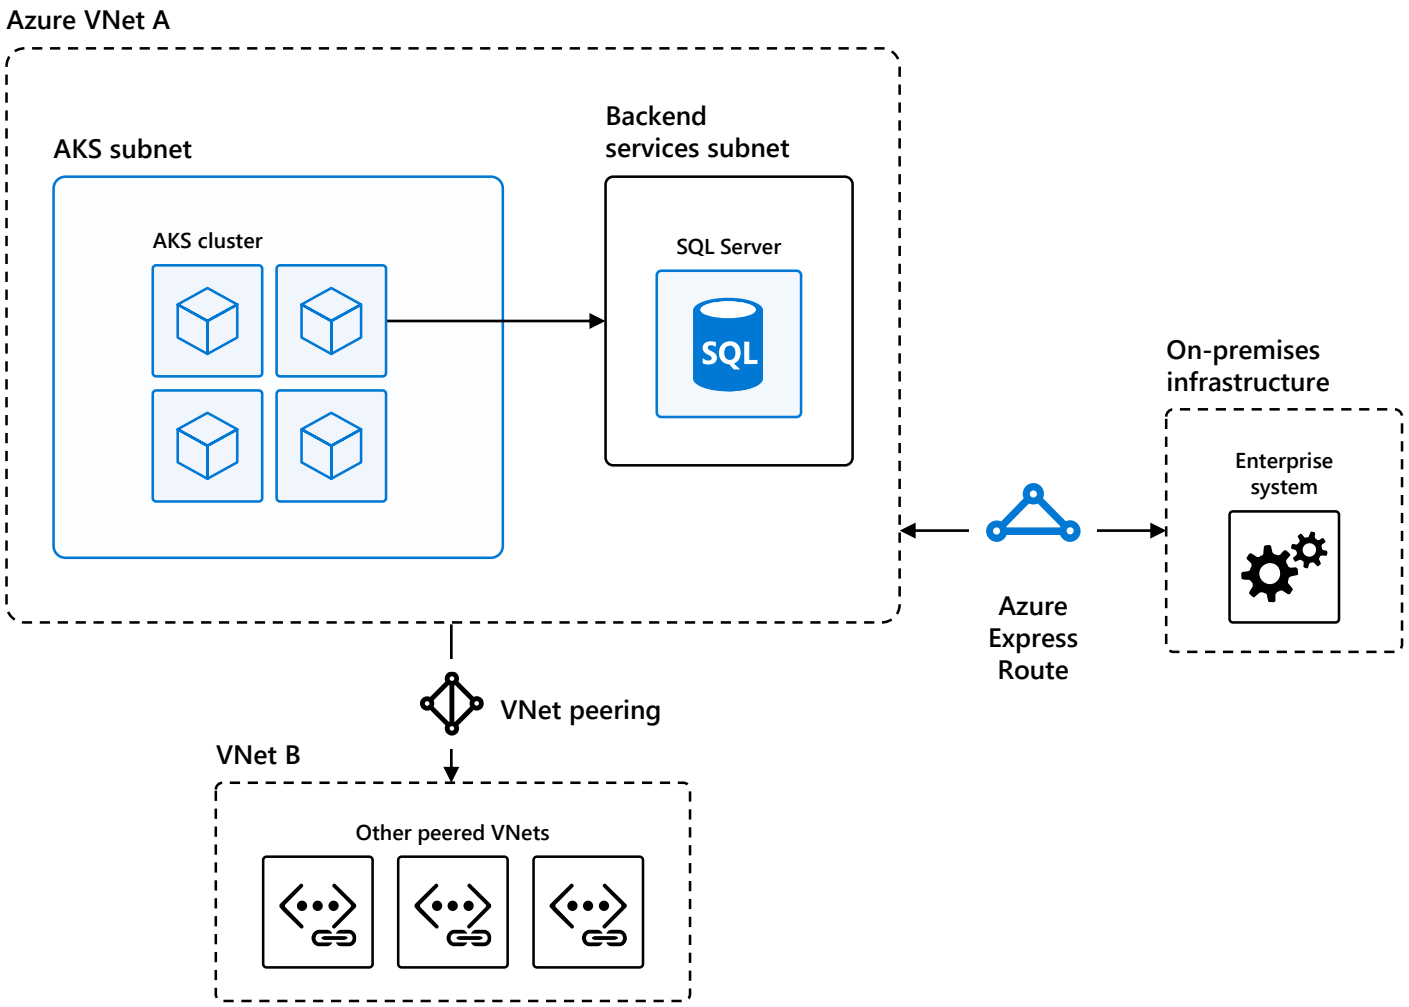

Public Service

- Service Type LoadBalancer
- Basic Layer4 Load Balancing (TCP/UDP)
- Each service as assigned an IP on the ALB

```
apiVersion: v1
kind: Service
metadata:
  name: frontendservice
spec:
  loadBalancerIP: X.X.X.X
  type: LoadBalancer
  ports:
  - port: 80
  selector:
    app: frontend
```

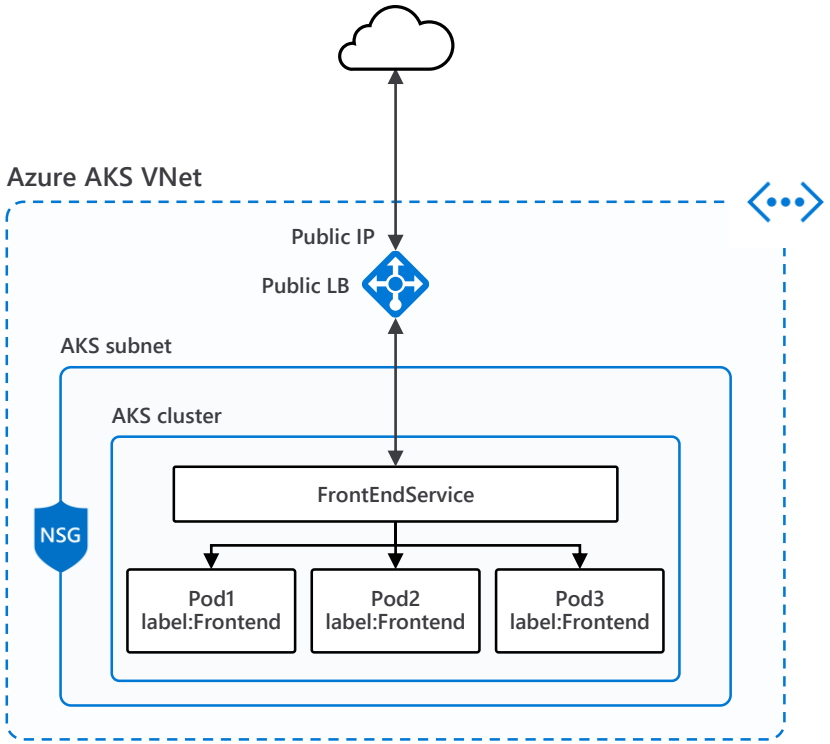

Internal Service

- Used for internal services that should be accessed by other VNETs or On-Premise only

```
apiVersion: v1
kind: Service
metadata:
  name: internalservice
  annotations:
    service.beta.kubernetes.io/azure-load-balancer-internal:
"true"
spec:
  type: LoadBalancer
  loadBalancerIP: 10.240.0.25
  ports:
  - port: 80
  selector:
   app: internal
```

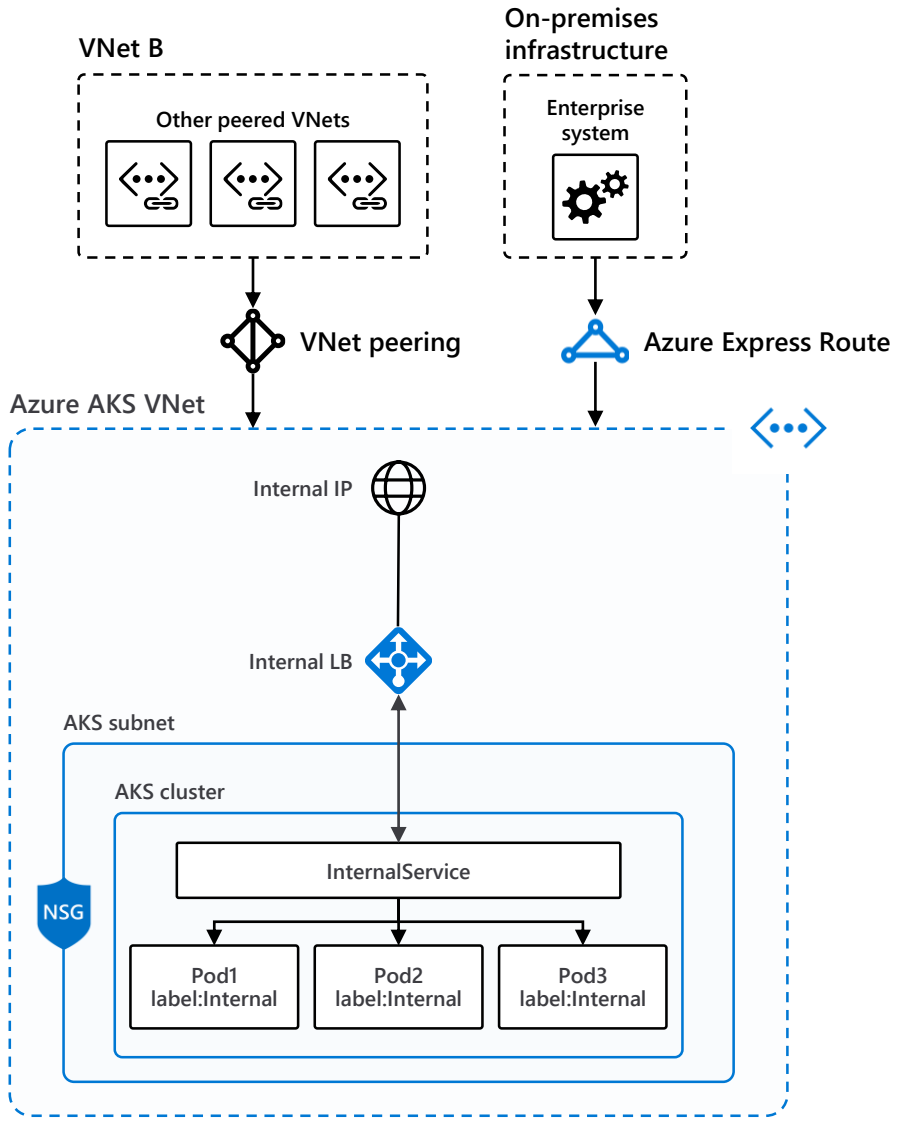

Ingress and Ingress Controllers

- **Ingress** is a Kubernetes API that manages external access to the services in the cluster
  - Supports HTTP and HTTPs
  - Path and Subdomain based routing
  - SSL Termination
  - Save on public Ips

- **Ingress controller** is a daemon, deployed as a Kubernetes Pod, that watches the Ingress Endpoint for updates. Its job is to satisfy requests for ingresses. Most popular one being **Nginx**.

Ingress

```
kind: Ingress
metadata:
 name: contoso-ingress
   annotations: kubernetes.io/ingress.class:
”PublicIngress"
spec:
 tls:
 - hosts:
   - contoso.com
   secretName: contoso-secret
 rules:
   - host: contoso.com
     http:
      paths:
      - path: /a
        backend:
         serviceName: servicea
         servicePort: 80
      - path: /b
        backend:
         serviceName: serviceb
         servicePort: 80
```

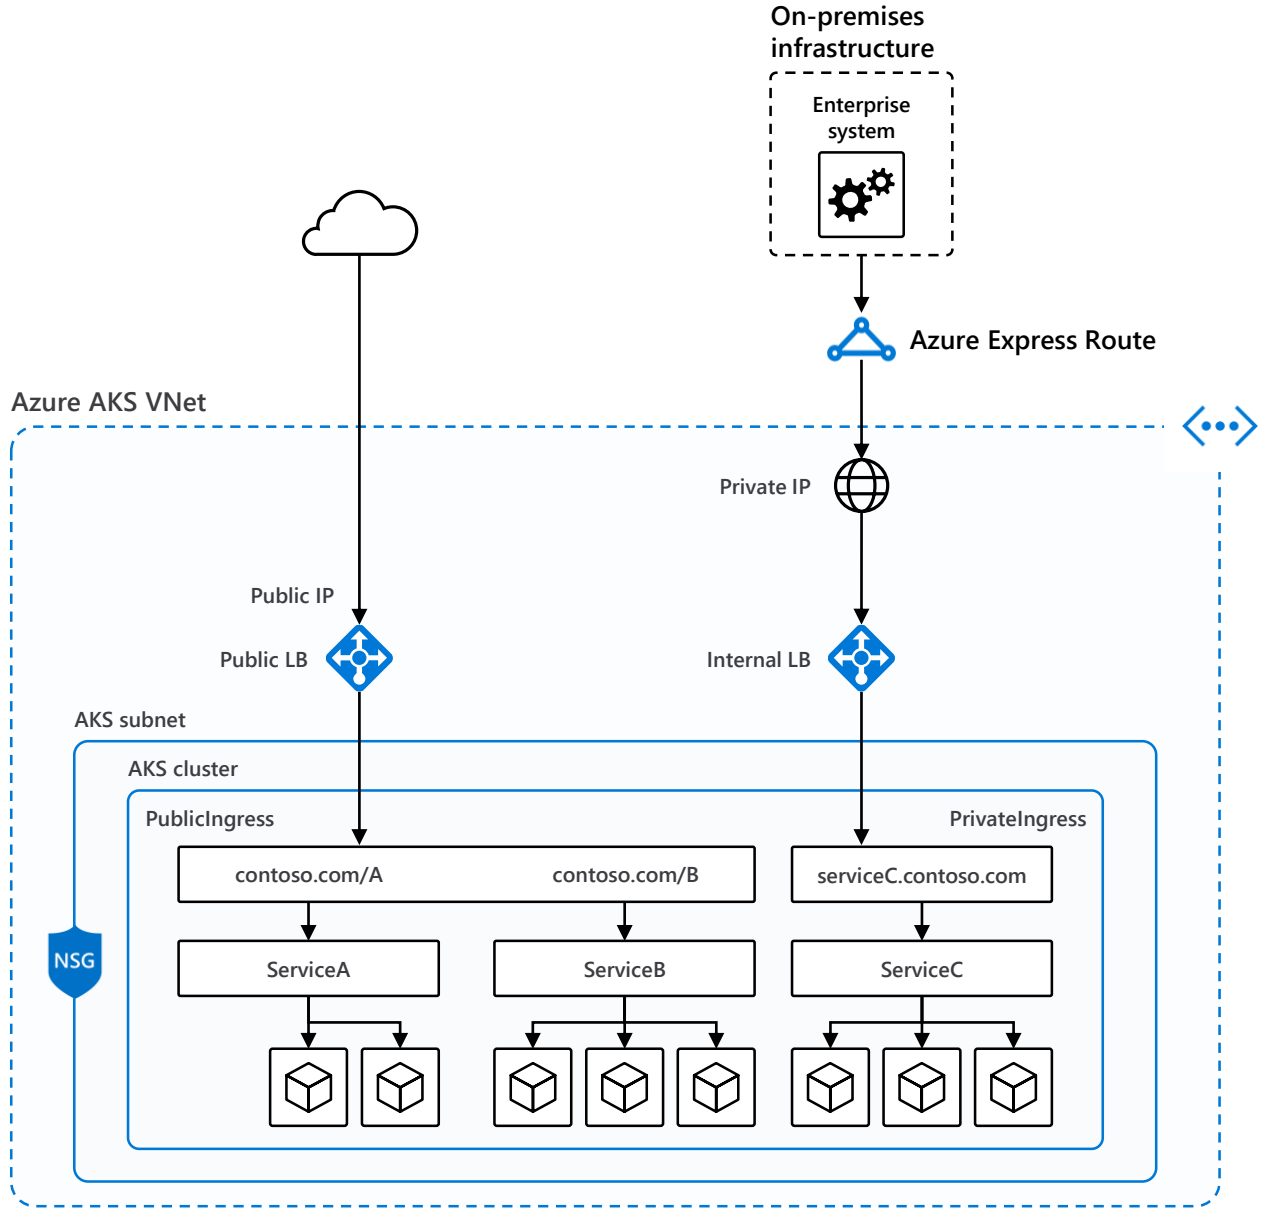

Securing Kubernetes Services with WAF

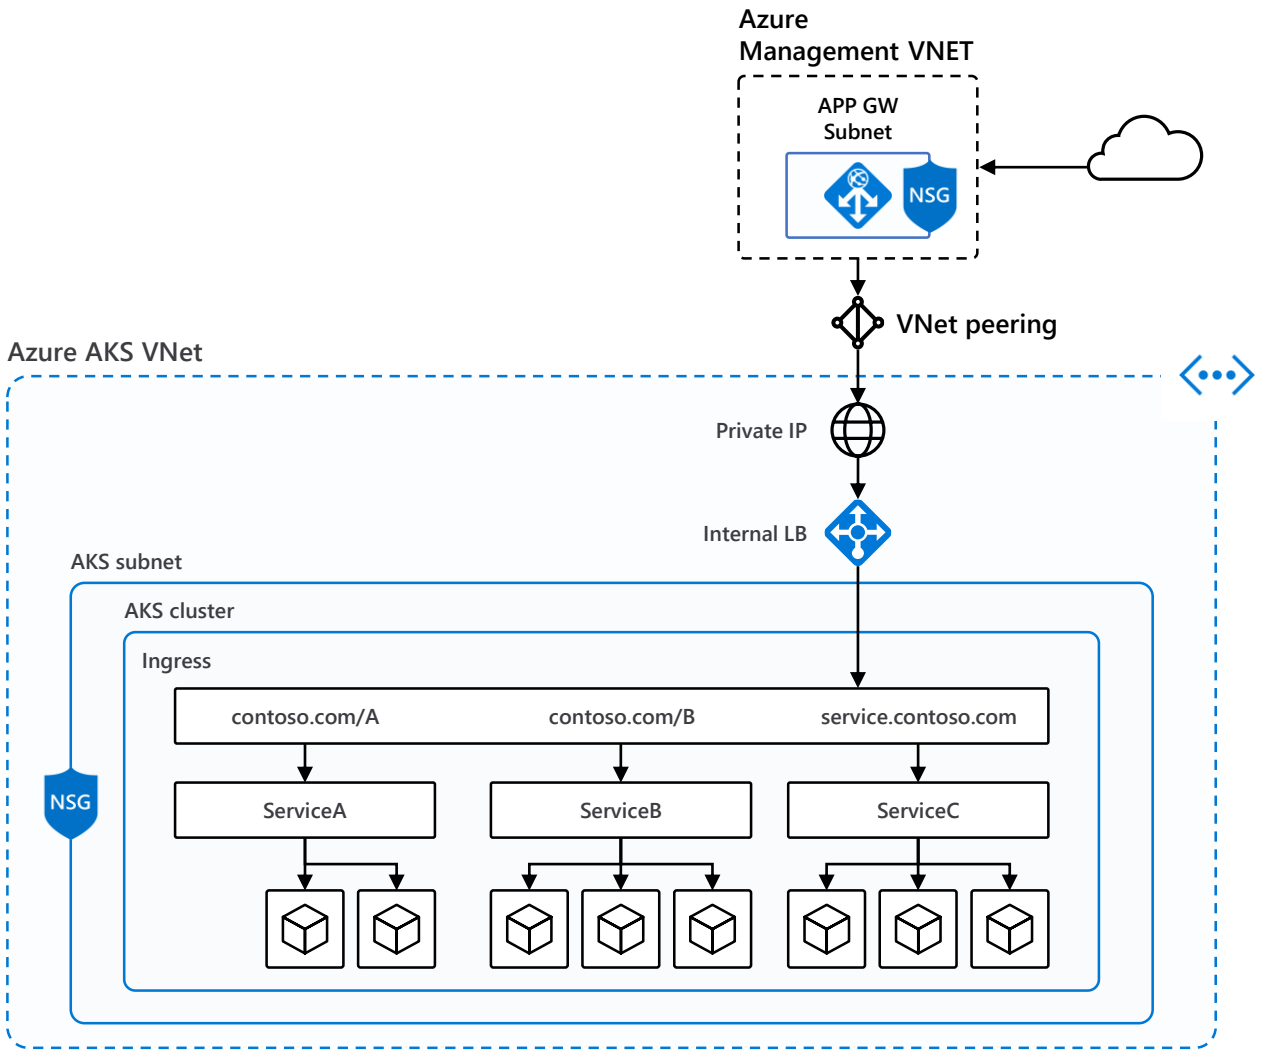

Cluster Management Through Bastion Host

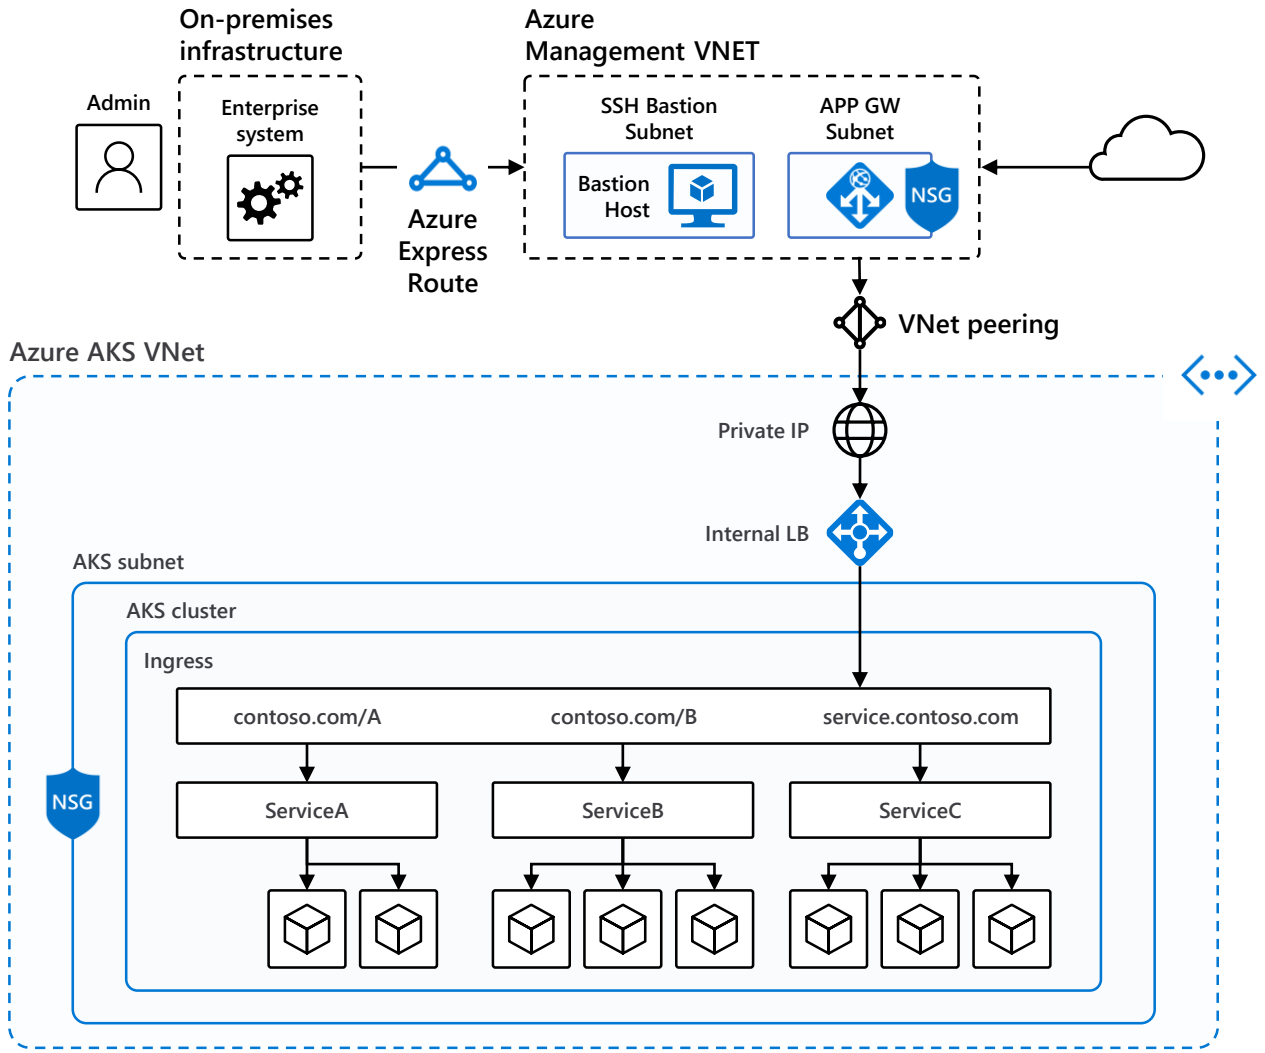

Demo

Ingress
Café Application with 2 endpoints Coffee and Tea
SSL termination
Loadbalancing

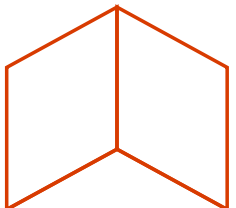

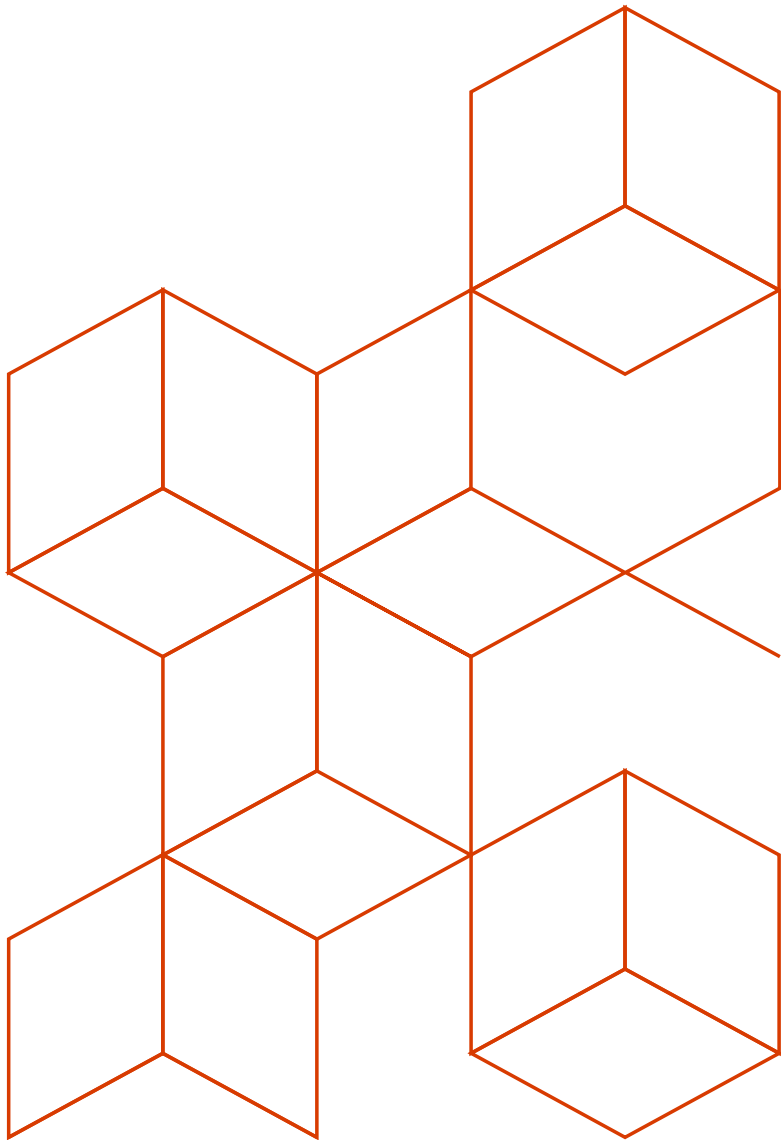

Summary

- Use AKS Advanced networking for seamless integration with your VNET
- Use Ingress and Ingress controllers for HTTP and HTTPs services
- Use Azure Application Gateway or any other alternative from the Azure Market place to secure your services using a WAF
- Use Bastion Hosts to access your nodes when needed

Securing your environment

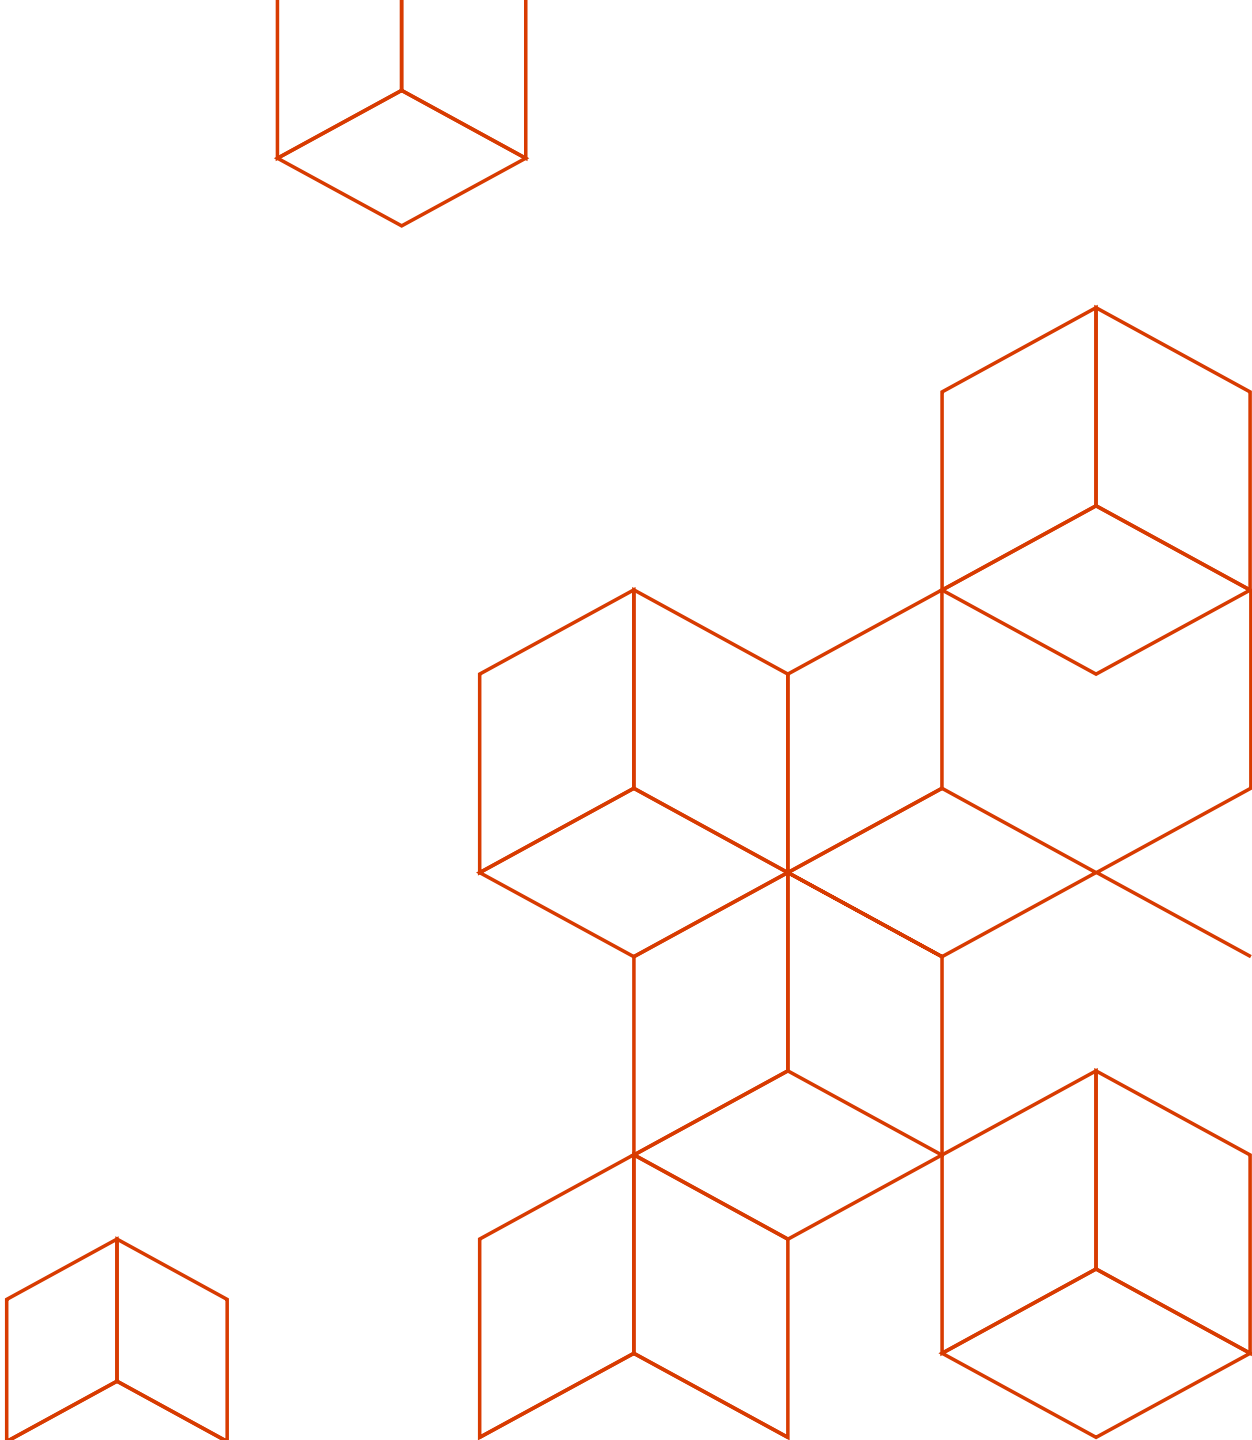

Cluster Level Security

Cluster Level Security

- Securing endpoints for API server and cluster nodes
  - Ensuring authentication and authorization (AAD + RBAC)
  - Setting up & keeping least privileged access for common tasks

Cluster Level - Identity and Access Management through AAD and RBAC

1. Kubernetes Developer authenticates with AAD
2. The AAD token issuance endpoint issues the access token
3. Developer performs action w/ AAD token.
   Eg. *kubectl create pod*
4. Kubernetes validates token with AAD and fetches the Developer's AAD Groups
   Eg. Dev Team A, App Group B
5. Kubernetes RBAC and cluster policies are applied
6. Request is successful or not based on the previous validation

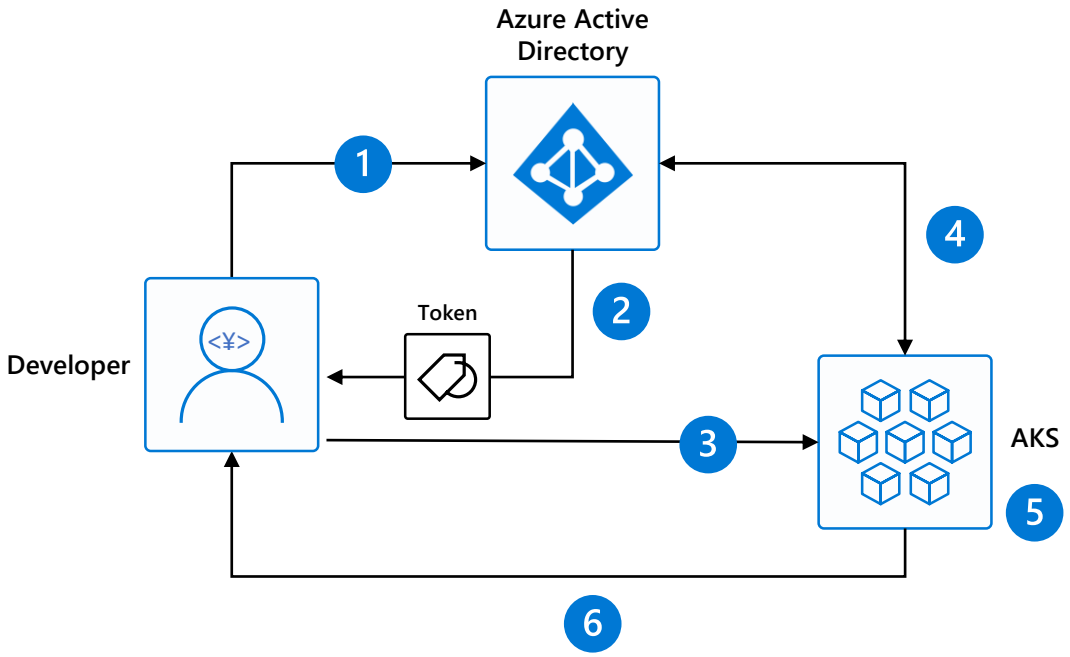

Provisioning AD-enabled AKS

```
$ az aks create --resource-group myAKSCluster --name myAKSCluster --generate-ssh-keys ¥
  --aad-server-app-id <Azure AD Server App ID> ¥
  --aad-server-app-secret <Azure AD Server App Secret> ¥
  --aad-client-app-id <Azure AD Client App ID> ¥
  --aad-tenant-id <Azure AD Tenant>
```

```
$ az aks get-credentials --resource-group myAKSCluster –name  myAKSCluster --admin
Merged "myCluster" as current context ..
```

```
$ kubectl get nodes

NAME                        STATUS    ROLES     AGE       VERSION
aks-nodepool1-42032720-0    Ready     agent     1h        v1.9.6
aks-nodepool1-42032720-1    Ready     agent     1h        v1.9.6
aks-nodepool1-42032720-2    Ready     agent     1h        v1.9.6
```

Provisioning AD-enabled AKS

Setting up a Cluster Role

```
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRole
metadata:
  labels:
   kubernetes.io/cluster-service: "true"
  name: cluster-admin
rules:
- apiGroups:
  - extensions
  - apps
  resources:
  - deployments
  verbs:
  - get
  - list
  - watch
  - update
  - patch
- apiGroups:
  - ""
  resources:
  - events
  - namespaces
  - nodes
  - pods
  verbs:
  - get
  - list
  - watch
```

Bind the Cluster Role to a user

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
name: contoso-cluster-admins
roleRef:
apiGroup: rbac.authorization.k8s.io
kind: ClusterRole
name: cluster-admin
subjects:
- apiGroup: rbac.authorization.k8s.io
kind: User
name: "user@contoso.com"
```

Bind the Cluster Role to a group

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
name: contoso-cluster-admins
roleRef:
apiGroup: rbac.authorization.k8s.io
kind: ClusterRole
name: cluster-admin
subjects:
- apiGroup: rbac.authorization.k8s.io
kind: Group
name: "894656e1-39f8-4bfe-b16a-510f61af6f41"
```

Azure Level - Identity and Access Management through AAD and RBAC

1. Kubernetes Administrator authenticates with AAD
2. The AAD token issuance endpoint issues the access token
3. Administrator fetches the admin kubeconfig and configures RBAC roles and bindings
4. Kubernetes Developer fetches the user kubeconfig

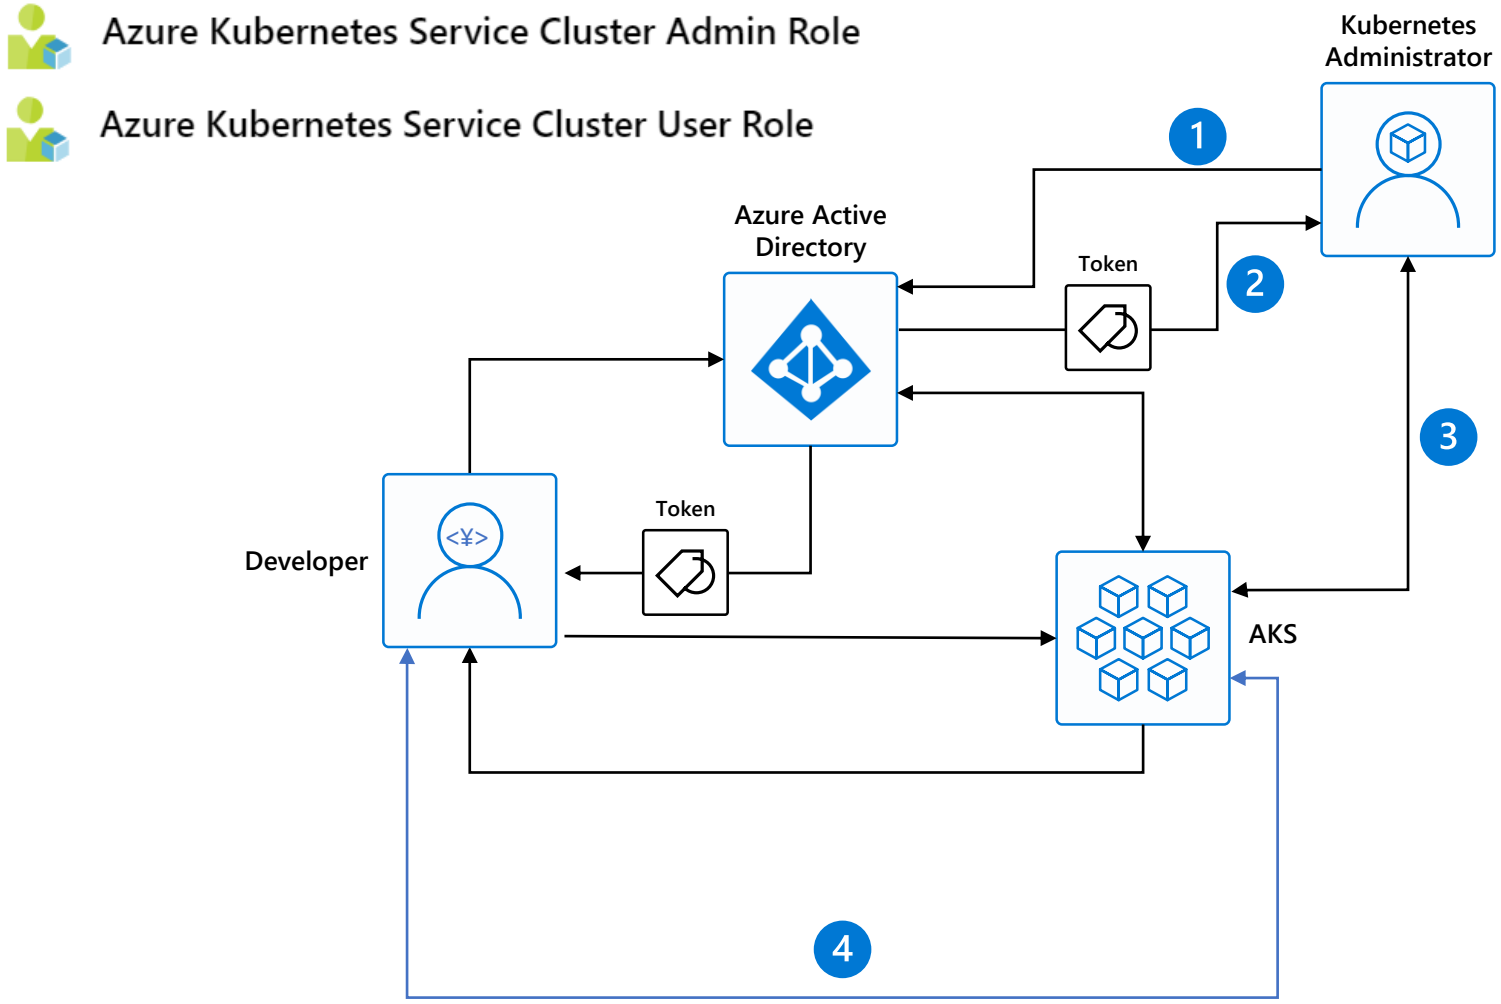

Provisioning AD-enabled AKS

```
$ az aks get-credentials --resource-group myAKSCluster --name myAKSCluster

$ kubectl get nodes

To sign in, use a web browser to open the page https://microsoft.com/devicelogin and
enter the code BUJHWDGNL to authenticate.

NAME                       STATUS    ROLES     AGE       VERSION
aks-nodepool1-42032720-0   Ready     agent     1h        v1.9.6
aks-nodepool1-42032720-1   Ready     agent     1h        v1.9.6
aks-nodepool1-42032720-2   Ready     agent     1h        v1.9.6
```

Or

```
Error from server (Forbidden): nodes is forbidden: User baduser@contoso.com cannot
list nodes at the cluster scope
```

Cluster Level Security

- Securing endpoints for API server and cluster nodes
  - Ensuring authentication and authorization (AAD + RBAC)
  - Setting up & keeping least privileged access for common tasks
  - Admission Controllers
    - DenyEscalatingExec
    - ValidatingAdmissionWebhooks
    - MutatingAdmissionWebhooks
    - ServiceAccount
    - Coming soon:
      - NodeRestriction
      - PodSecurityPolicy

Cluster Level – Nodes, Upgrade and Patches

- Regular maintenance, security and cleanup tasks
  - Maintain, update and upgrade hosts and kubernetes
  - Monthly ideal, 3 months minimum
  - Security patches
    - AKS automatically applies security patches to the nodes on a nightly schedule
    - You're responsible to reboot as required
    - Kured DaemonSet https://github.com/weaveworks/kured

Upgrade to version 1.10.6

```
$ az aks upgrade --name myAKSCluster ¥
--resource-group myResourceGroup ¥
--kubernetes-version 1.10.6
```

- **SSH Access**
  - DenyEscalatingExec
- **Running benchmarks and tests to validate cluster setup**
  - Kube-bench
  - Aqua Hunter
  - Others

Container Level Security and Isolation

Container Level – The images

- Trusted Registry
- Regularly apply security updates to the container images

Container Level – Images and Runtime

- Scan your images, scan your containers
- Runtime enforcement and remediation

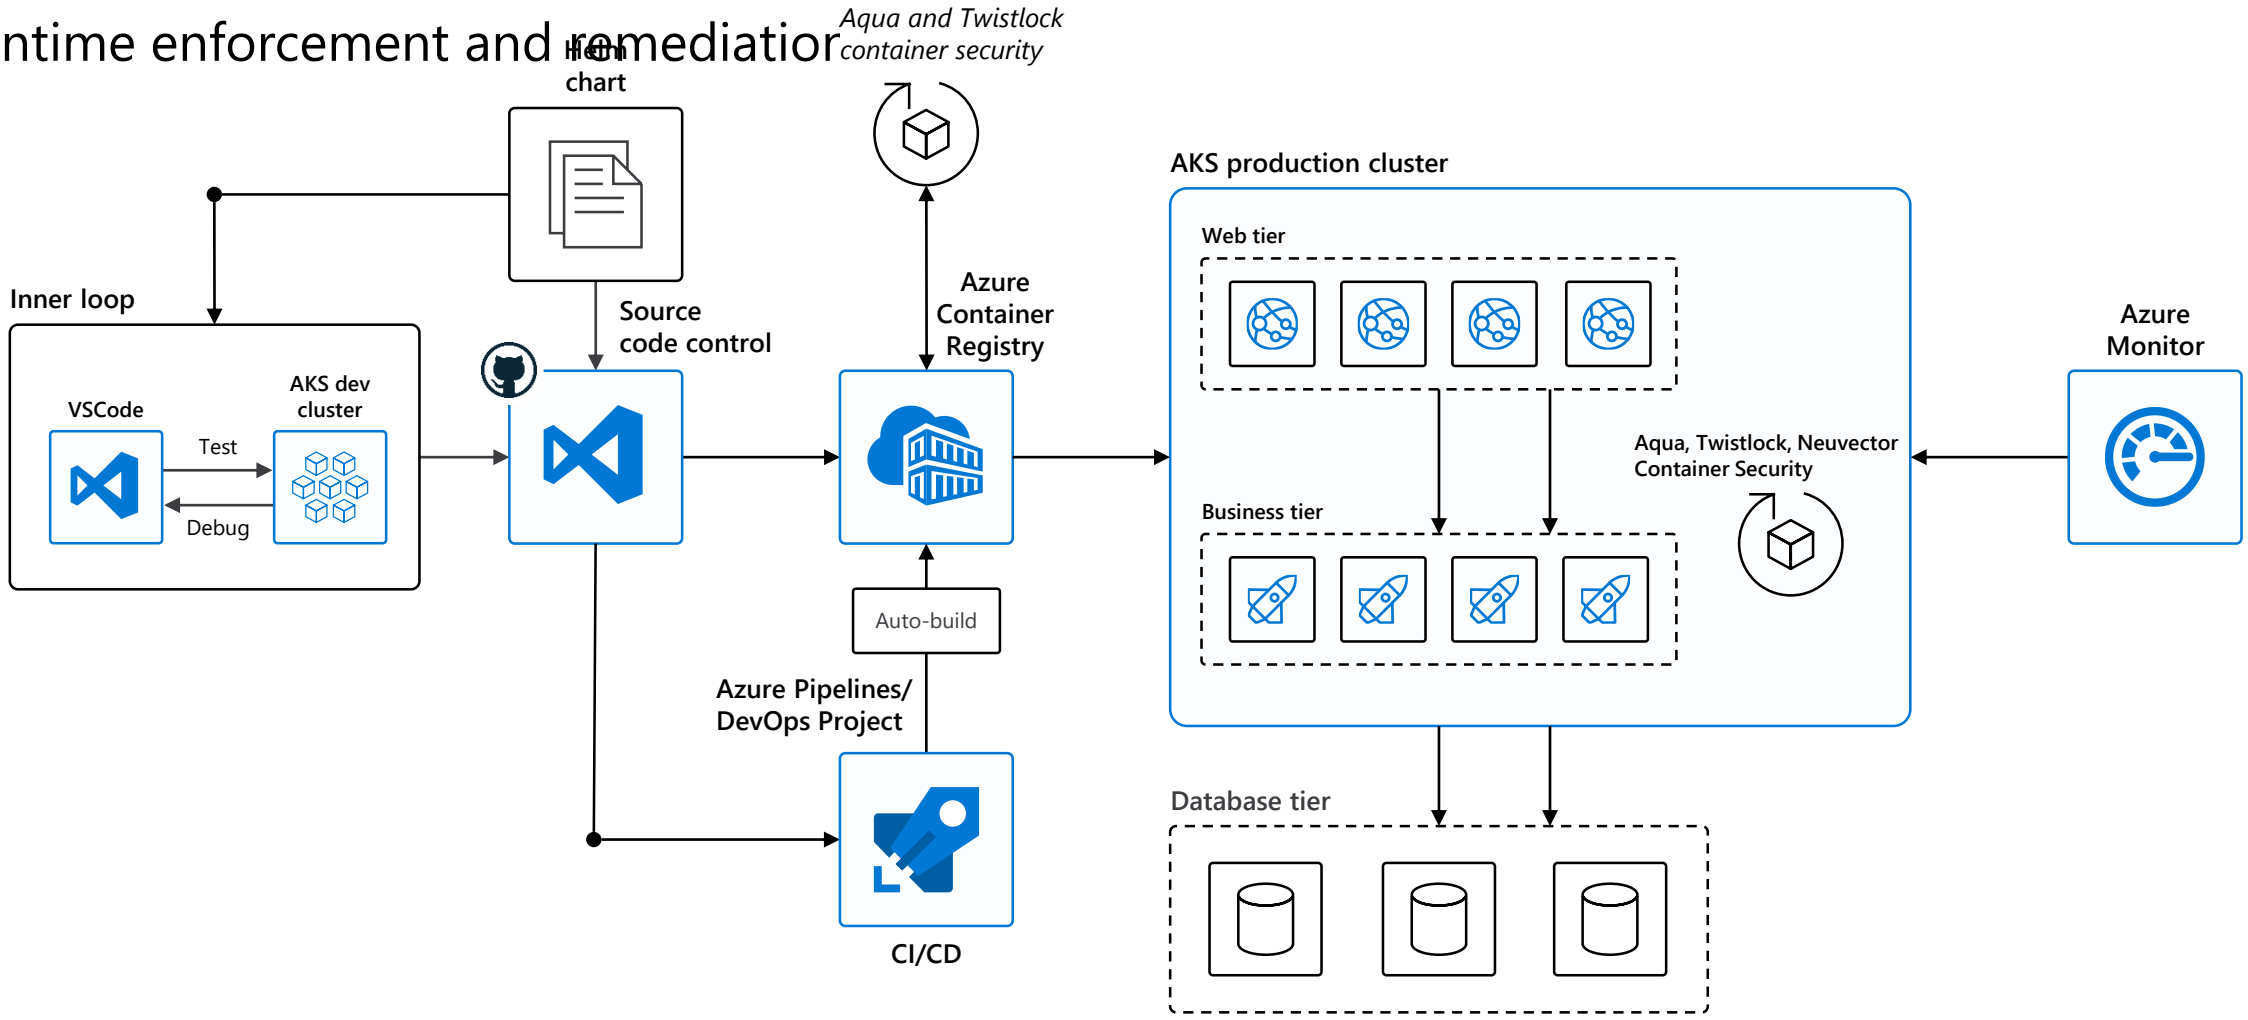

Container Level – The access

- Avoid access to HOST IPC namespace - only if absolutely necessary
- Avoid access to Host PID namespace - only if absolutely necessary
- Avoid root / privileged access
    - Consider Linux Capabilities

Container Level – App Armor

Securing a Pod with a deny-write.profile

```
apiVersion: v1
kind: Pod
metadata:
  name: hello-apparmor
  annotations:
    container.apparmor.security.beta.kubernetes.io/hello: localhost/k8s-apparmor-example-deny-write
spec:
  containers:
  - name: hello
    image: busybox
    command: [ "sh", "-c", "echo 'Hello AppArmor!' && sleep 1h" ]
```

```
$ kubectl exec hello-apparmor touch /tmp/test
touch: /tmp/test: Permission denied
error: error executing remote command: command terminated with non-zero exit
code: Error executing in Docker Container: 1
```

deny-write.profile

```
#include <tunables/global>

profile k8s-apparmor-example-deny-write flags=(attach_disconnected) {
  #include <abstractions/base>

  file,

  # Deny all file writes.
  deny /** w,
}
```

Container Level - Seccomp

Securing a Pod with a prevent-chmod profile

```
apiVersion: v1
kind: Pod
metadata:
  name: chmod-prevented
  annotations:
    seccomp.security.alpha.kubernetes.io/pod:
localhost/prevent-chmod

spec:
  containers:
  - name: chmod
    image: busybox
    command:
      - "chmod“
    args:
      - “777“
      - /etc/hostname
  restartPolicy: Never
```

Seccomp Profile /var/lib/kubelet/seccomp/prevent-chmod

```
{
  "defaultAction": "SCMP_ACT_ALLOW",
  "syscalls": [
    {
      "name": "chmod",
      "action": "SCMP_ACT_ERRNO"
    }
  ]
}
```

Container Level

```
$ kubectl create -f seccomp-pod.yaml
pod "chmod-prevented" created

$ kubectl get pods
NAME                        READY     STATUS    RESTARTS   AGE
chmod-prevented             0/1       Error     0          8s
```

Pod Level Security

Pod Level – Pod Security Context

```yaml
apiVersion: v1
kind: Pod
metadata:
  name: security-context-demo
spec:
  securityContext:
    runAsUser: 1000
    fsGroup: 2000
  volumes:
  - name: sec-ctx-vol
    emptyDir: {}
  containers:
  - name: sec-ctx-demo
    image: ignite.azurecr.io/nginx-demo
    volumeMounts:
    - name: sec-ctx-vol
      mountPath: /data/demo
    securityContext:
      runAsUser: 2000
      allowPrivilegeEscalation: false
      capabilities:
        add: ["NET_ADMIN", "SYS_TIME"]
      seLinuxOptions:
        level: "s0:c123,c456"
```

Pod Level – Pod Security Context

```yaml
apiVersion: v1
kind: Pod
metadata:
  name: security-context-demo
spec:
  securityContext:
    runAsUser: 1000
    fsGroup: 2000
  volumes:
  - name: sec-ctx-vol
    emptyDir: {}
  containers:
  - name: sec-ctx-demo
    image: ignite.azurecr.io/nginx-demo
    volumeMounts:
    - name: sec-ctx-vol
      mountPath: /data/demo
    securityContext:
      runAsUser: 2000
      allowPrivilegeEscalation: false
      capabilities:
        add: ["NET_ADMIN", "SYS_TIME"]
      seLinuxOptions:
        level: "s0:c123,c456"
```

Pod Level – Pod Security Policies

```yaml
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: restricted
  annotations:
    seccomp.security.alpha.kubernetes.io/allowedProfileNames: 'docker/default’
    apparmor.security.beta.kubernetes.io/allowedProfileNames: 'runtime/default’
    seccomp.security.alpha.kubernetes.io/defaultProfileName: 'docker/default’
    apparmor.security.beta.kubernetes.io/defaultProfileName: 'runtime/default’
spec:
  privileged: false
  allowPrivilegeEscalation: false # Required to prevent escalations to root.
  requiredDropCapabilities: # This is redundant with non-root + disallow privilege escalation, but we can provide it for defense in depth.
    - ALL
  volumes: # Allow core volume types.
    - 'configMap’
    - 'emptyDir’
    - 'projected’
    - 'secret’
    - 'downwardAPI’
    - 'persistentVolumeClaim’ # Assume that persistentVolumes set up by the cluster admin are safe to use.
  hostNetwork: false
  hostIPC: false
  hostPID: false
  runAsUser:
    rule: 'MustRunAsNonRoot’ # Require the container to run without root privileges.
  seLinux:
    rule: 'RunAsAny’ # This policy assumes the nodes are using AppArmor rather than SELinux.
  supplementalGroups:
    rule: 'MustRunAs’
    ranges:
      - min: 1      # Forbid adding the root group.
        max: 65535
  fsGroup:
    rule: 'MustRunAs’
    ranges:
      - min: 1     # Forbid adding the root group.
        max: 65535
  readOnlyRootFilesystem: false
```

Pod level

- Pod Security Context
- Pod Security Policies
- AlwaysPull Images

Securing Workloads

Pod Identity

1. Kubernetes operator defines an identity map for K8s service accounts

2. Node Managed Identity (NMI) watches for mapping reaction and syncs to Managed Service Identify (MSI)

3. Developer creates a pod with a service account. Pod uses standard Azure SDK to fetch a token bound to MSI

4. Pod uses access token to consume other Azure services; services validate token

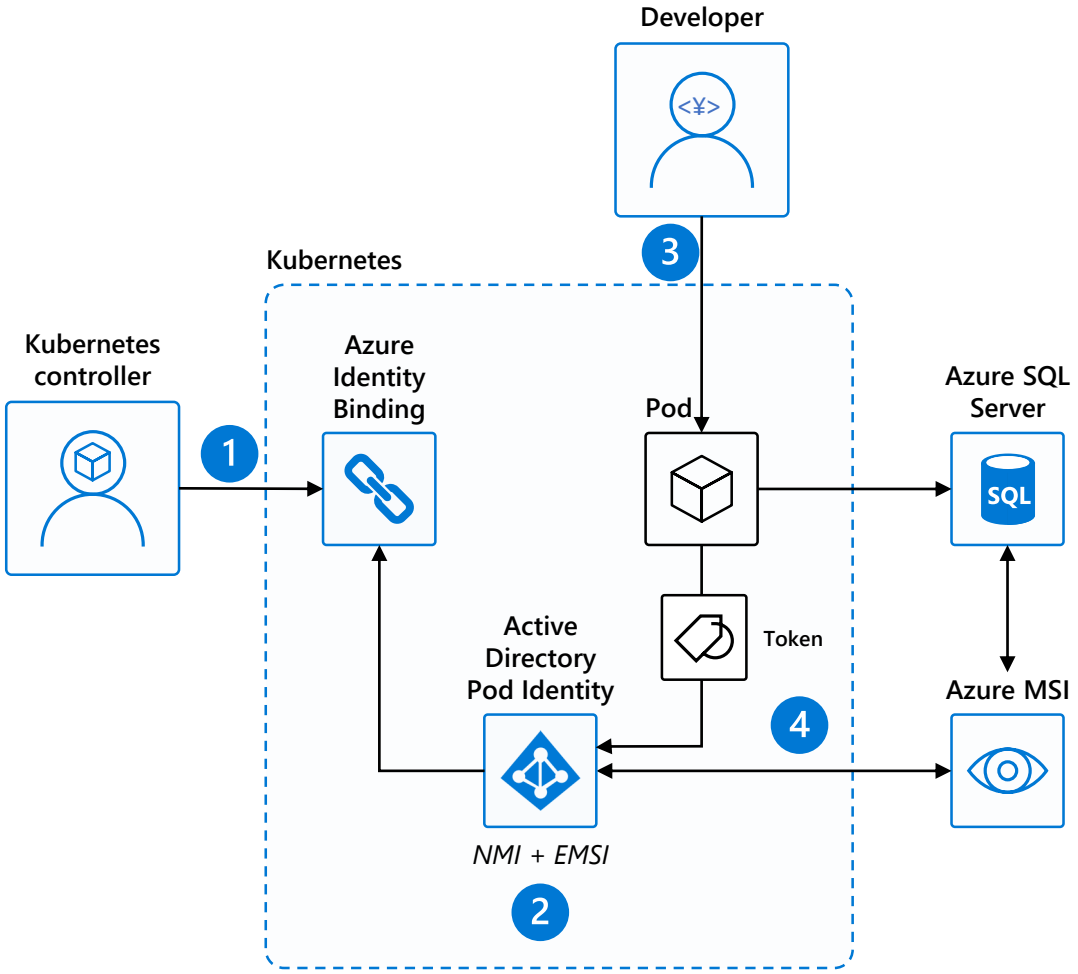

Securing workloads

- Managing secrets and privileged information
  - Azure Key Vault

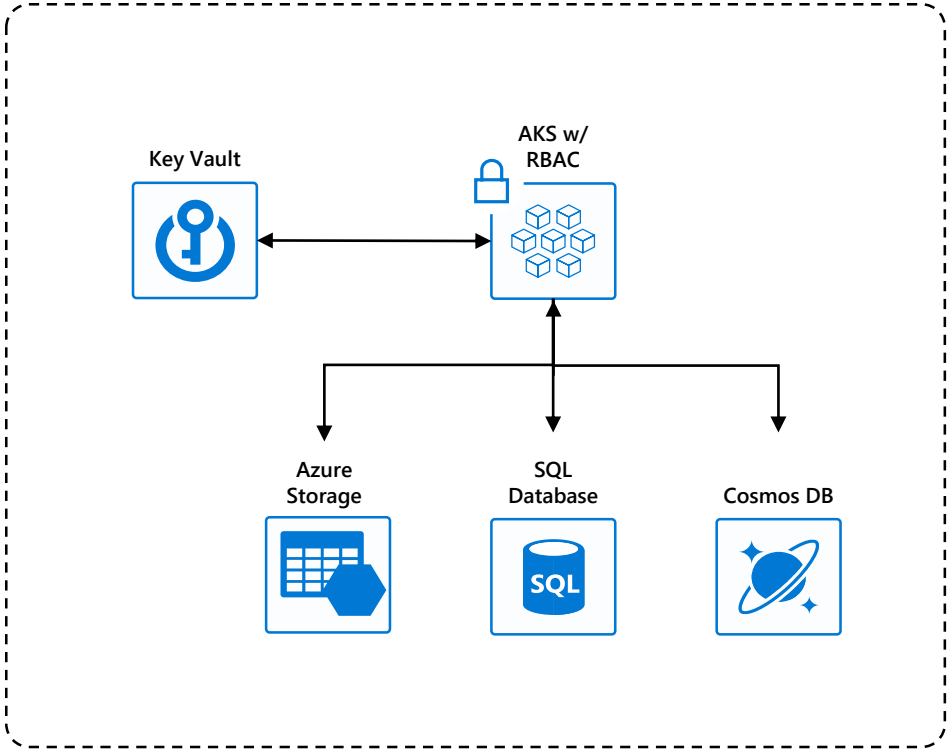

Demo

Protect your secrets with Azure Key Vault

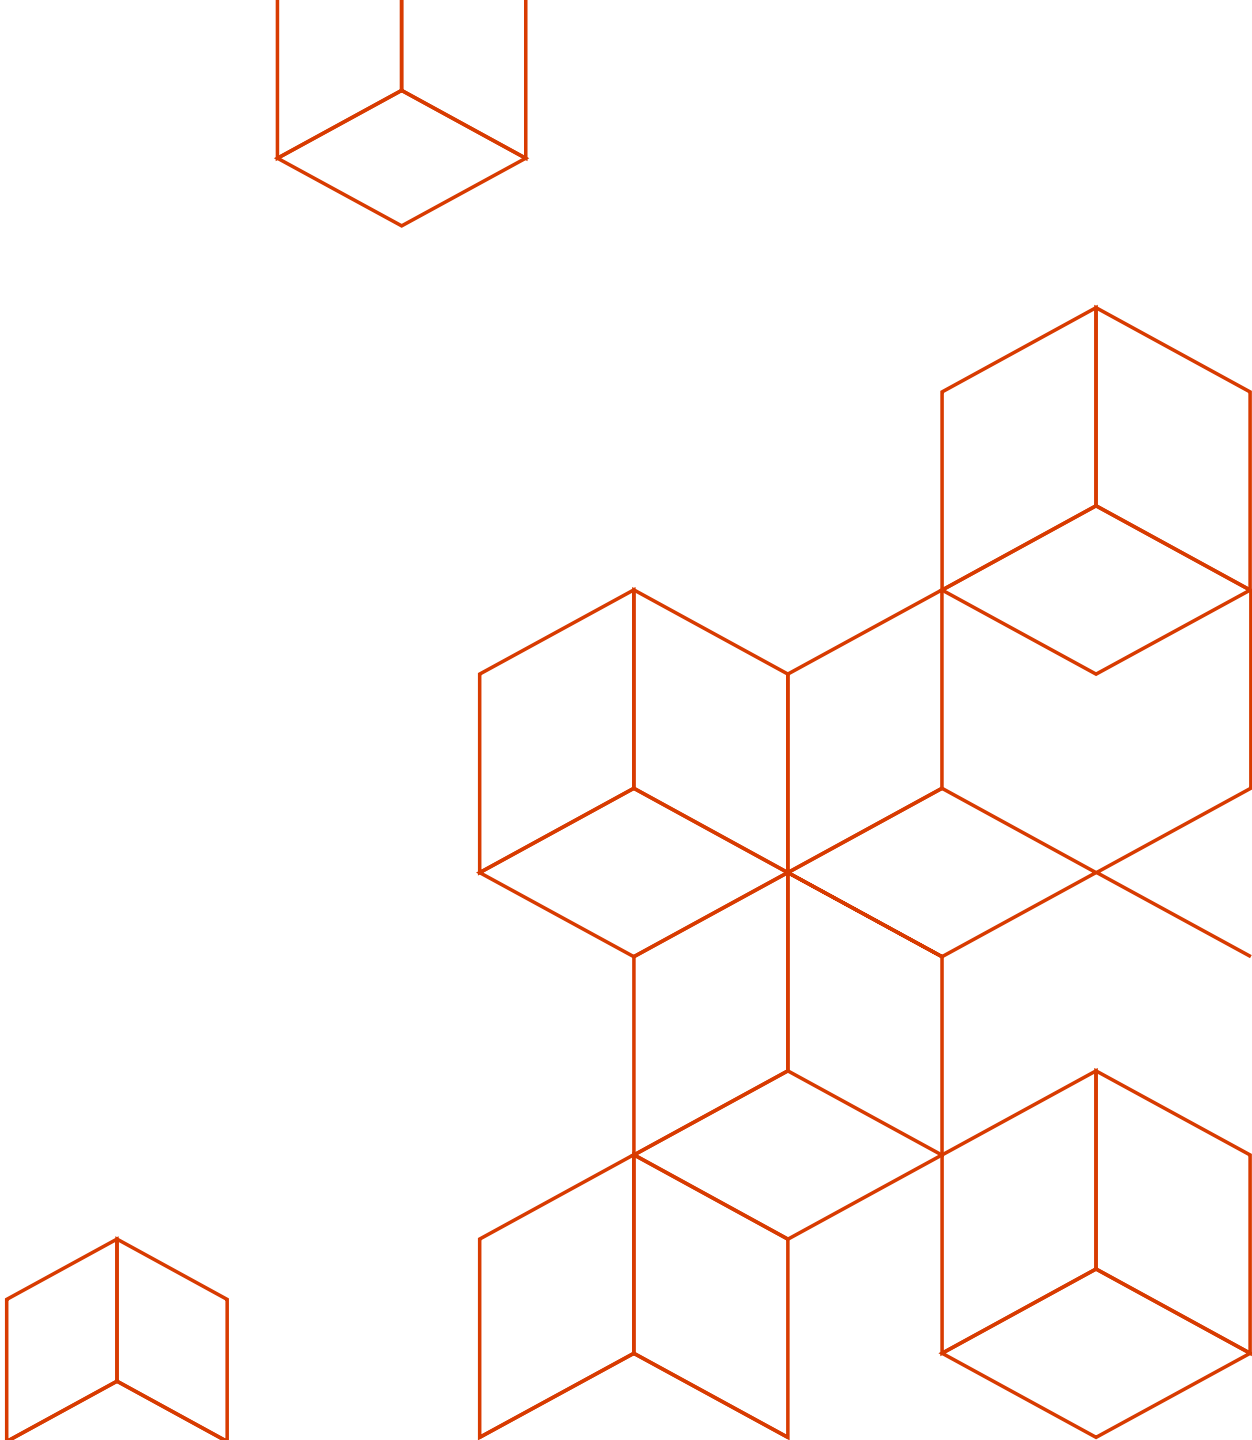

Securing workloads

- Service Endpoints
- Filter secrets from the logs
- Encrypted Service to Service Communication
  - mTLS between services
  - Service Meshes

Compliance

- AKS is SOC 1/2 , PCI , HIPPA and ISO certified
- All the details are listed in the Azure Trust Center

Autoscaling Applications and Clusters

Manual scaling is tedious and ineffective

- Horizontal pod autoscaling(HPA) -> Scaling pods/containers
- Cluster Autoscaling -> Scaling infrastructure/VM’s
- AKS + ACI + VK for burst scenarios -> Scaling pods/containers

How HPA works?

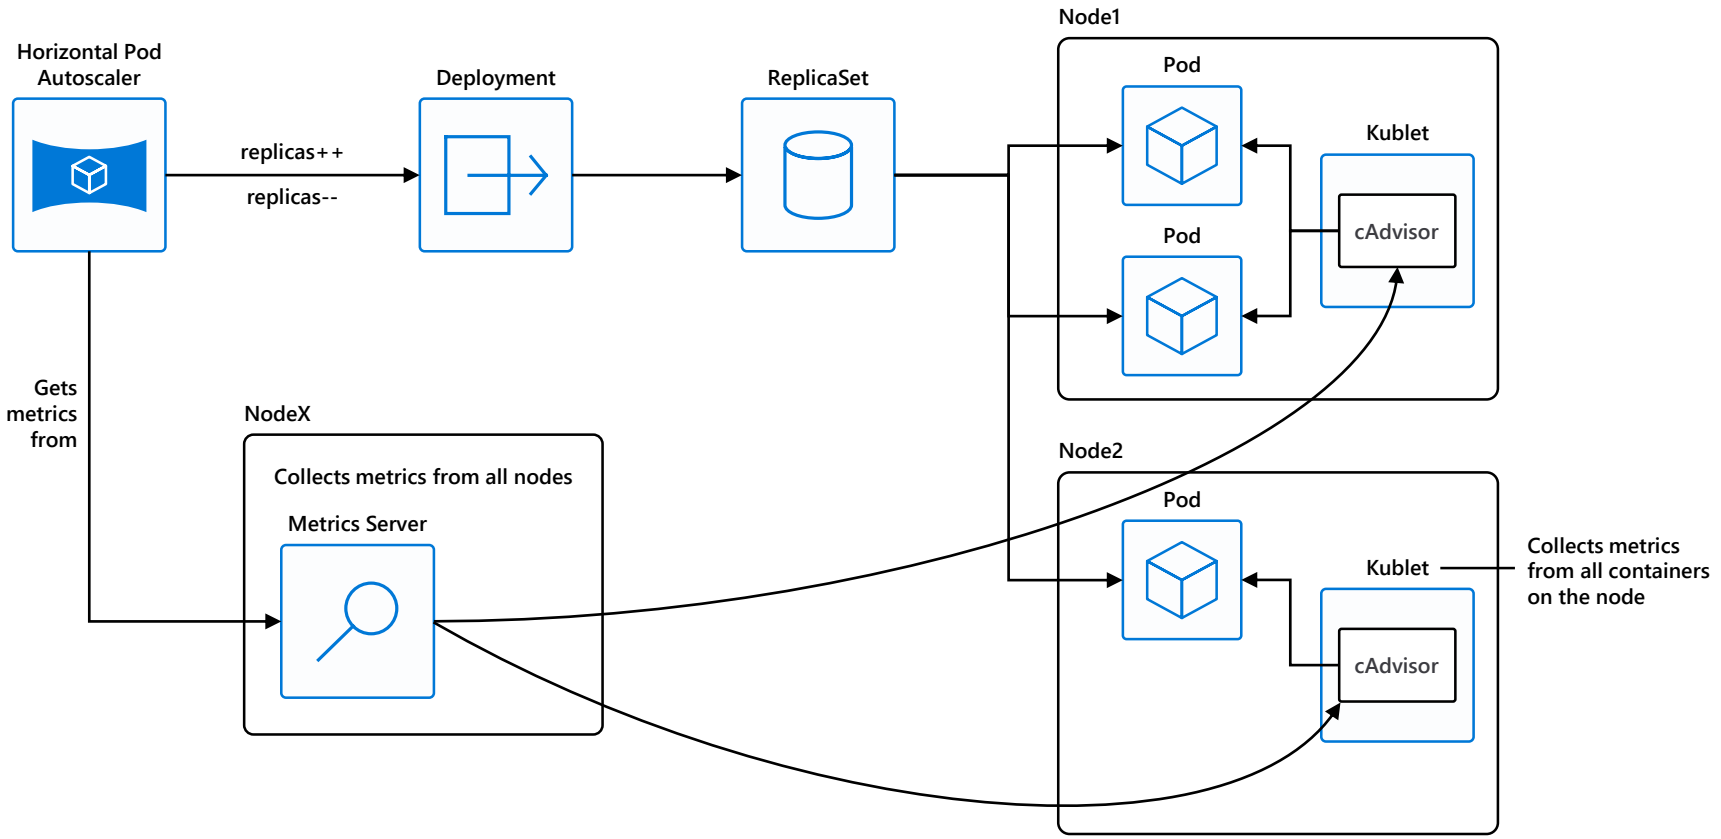

Cluster Autoscaler

- Scales nodes based on pending pods
- Scale up and scale down
- Reduces dependency on monitoring
- Removes need for users to manage nodes and monitor service usage manually

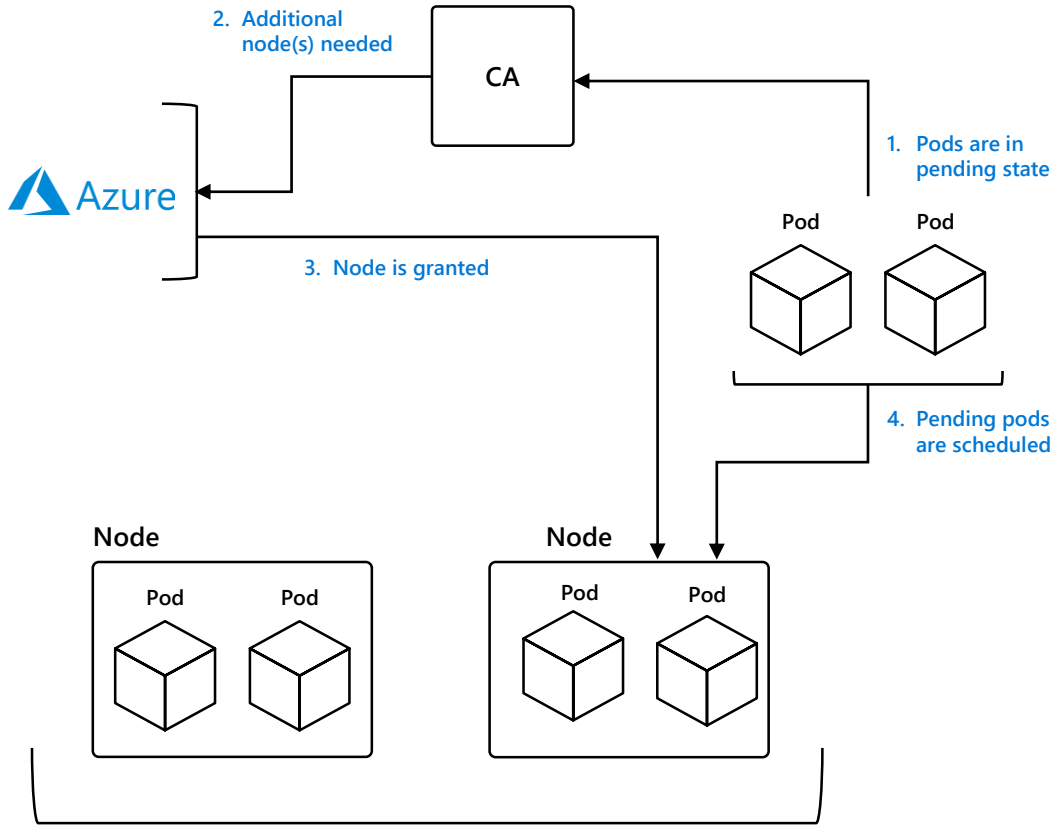

Bursting with the ACI Connector/ Virtual Kubelet

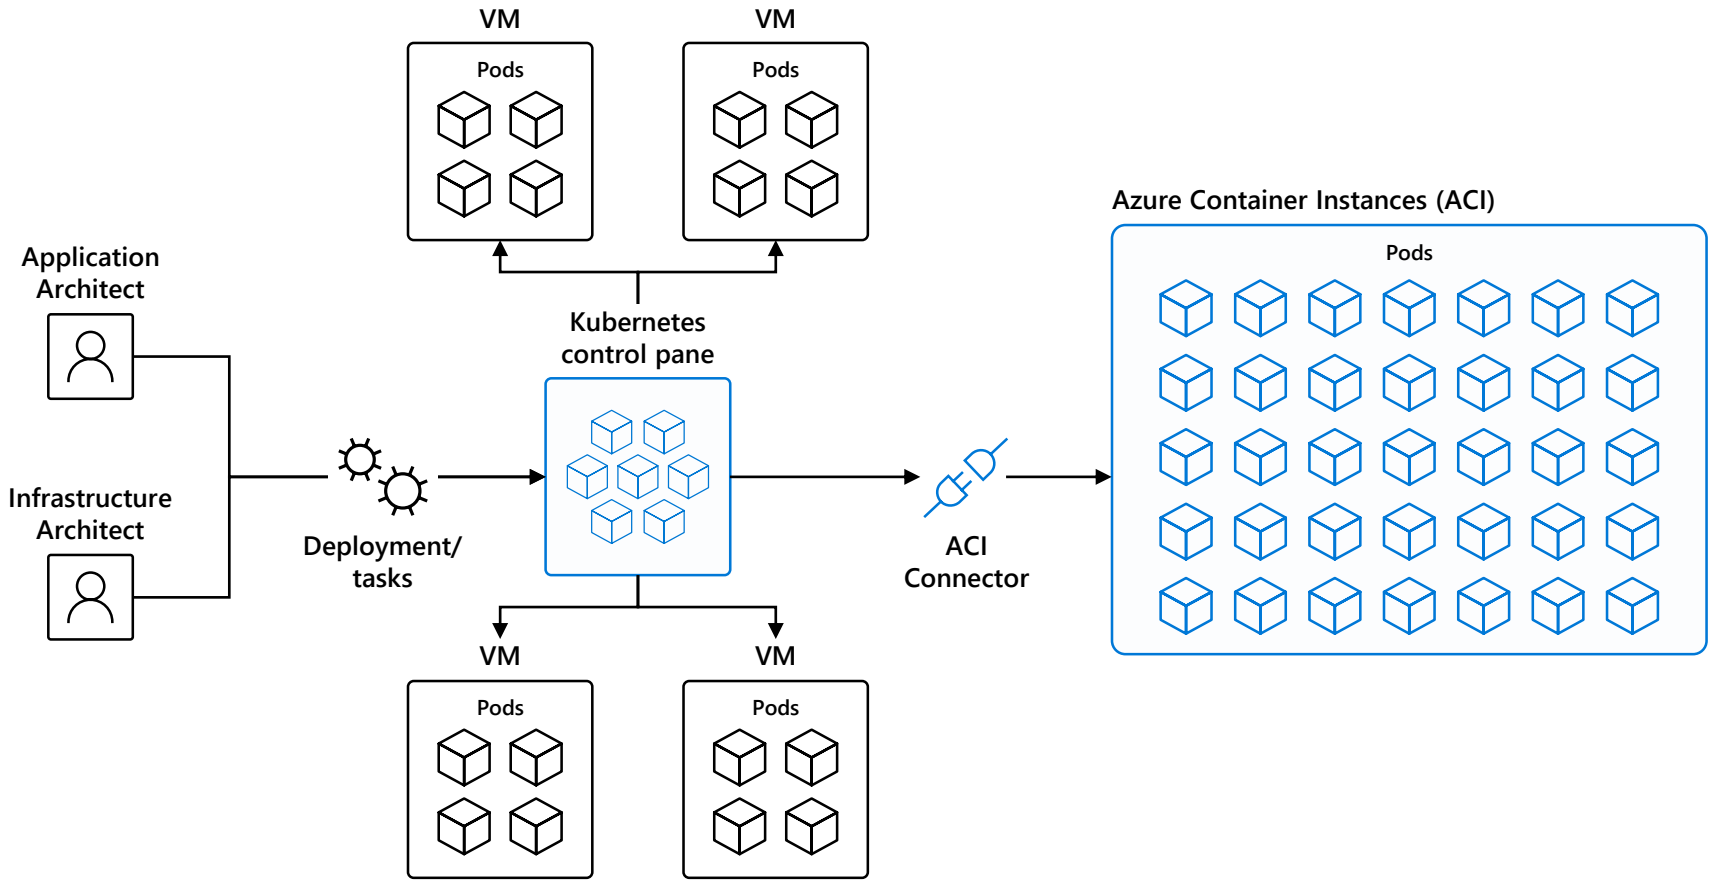

Fast container autoscaling

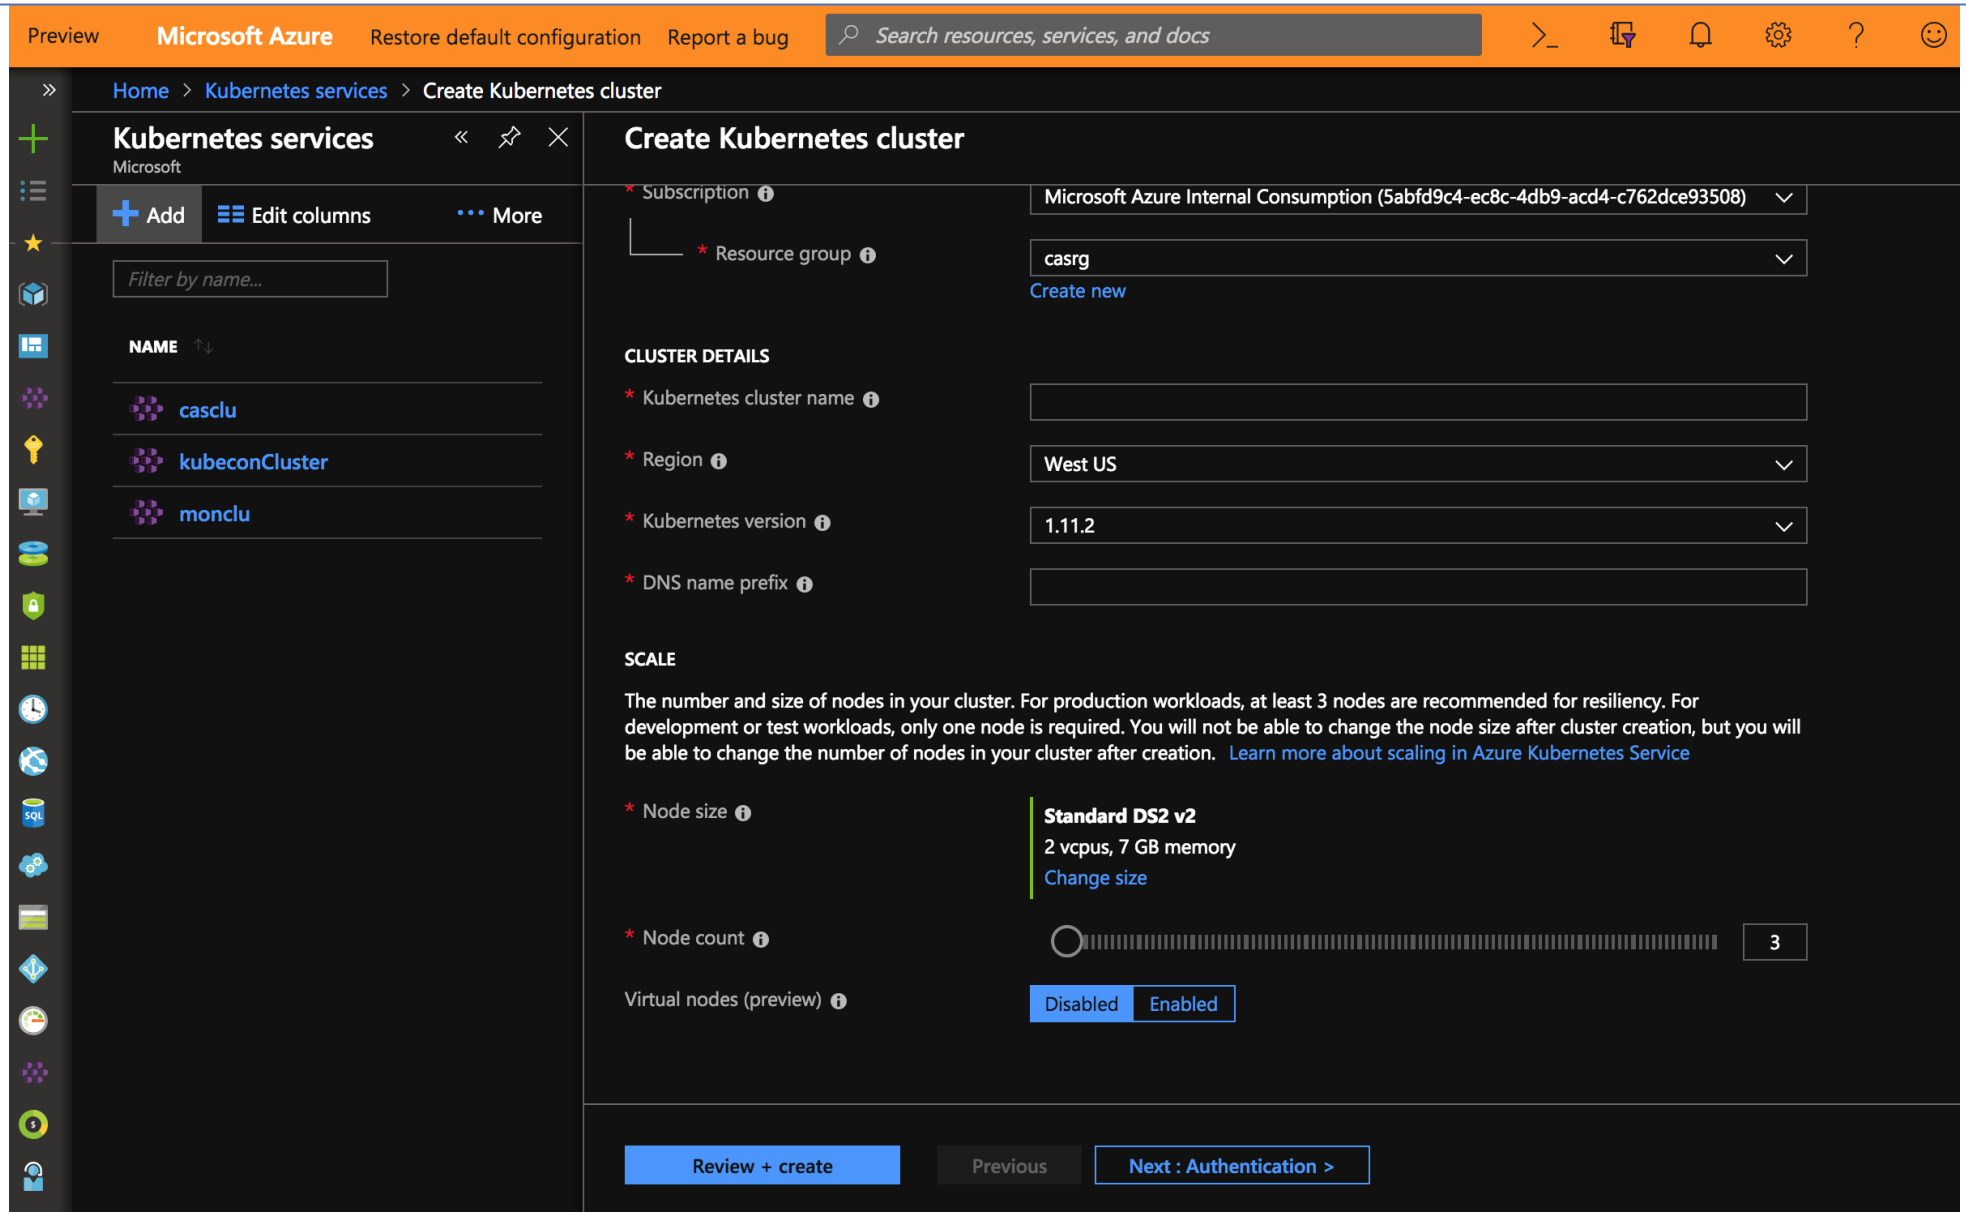

Demo

Cluster autoscaler
AKS + VK burst ACI

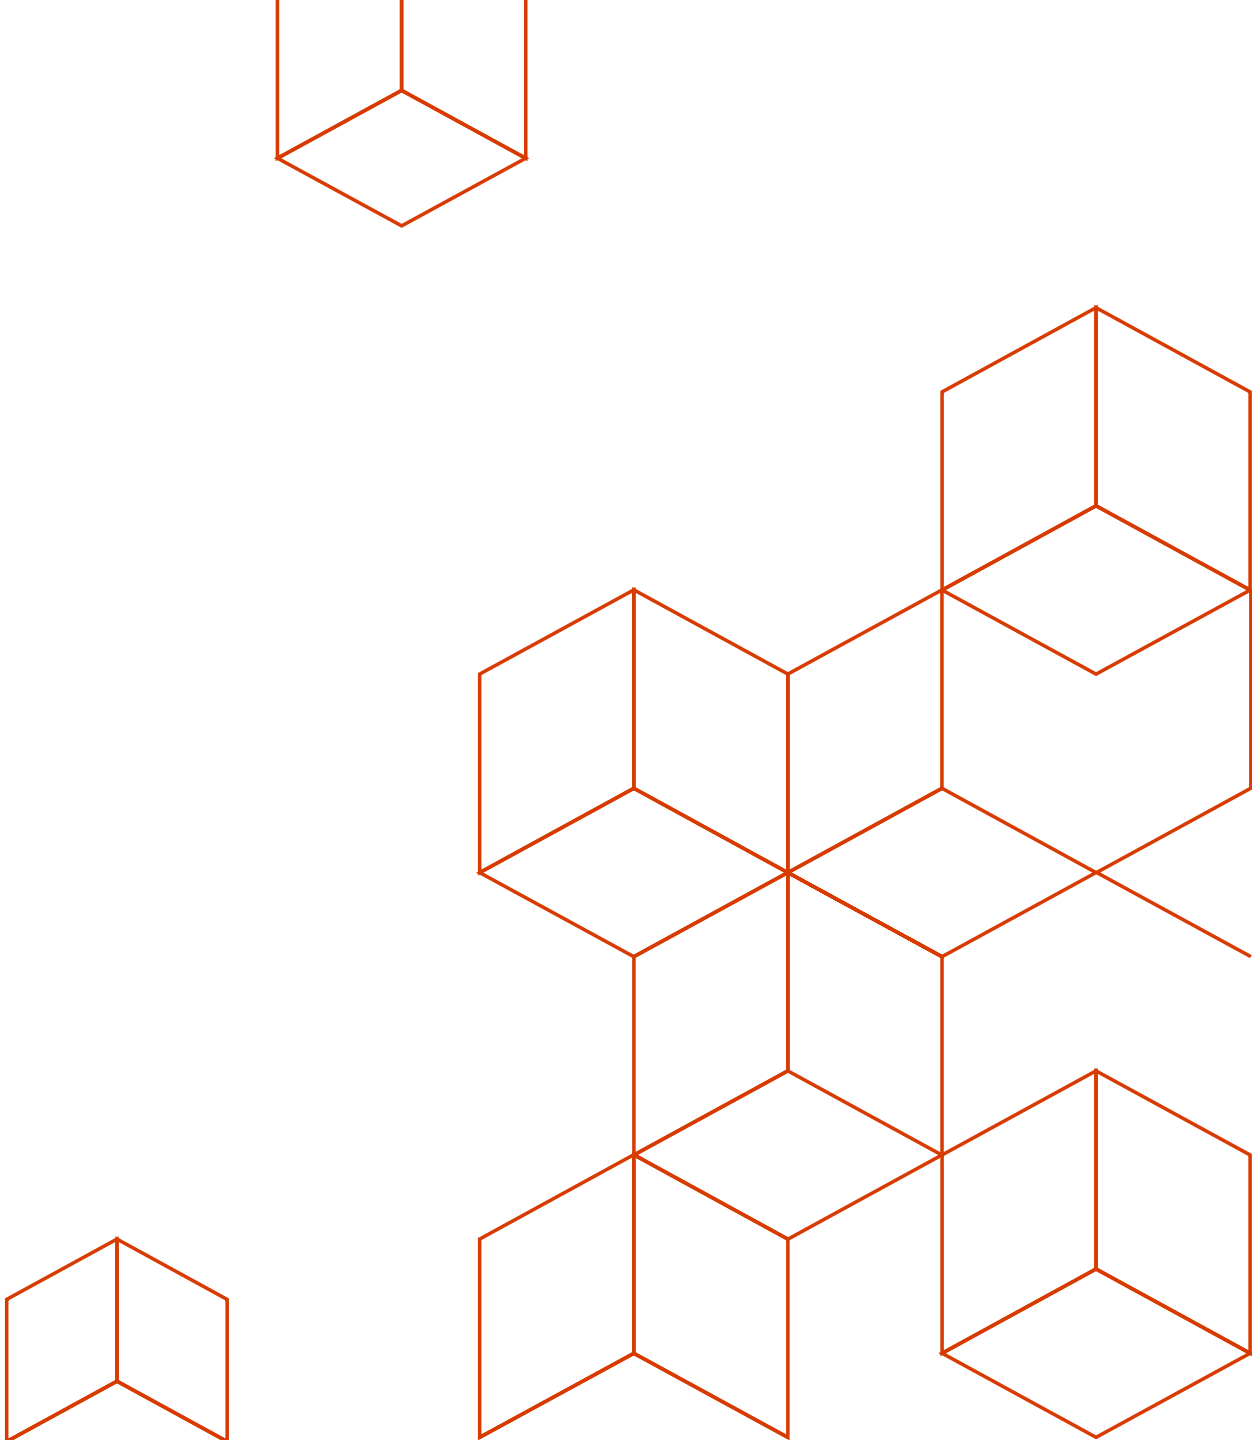

Multi-Region

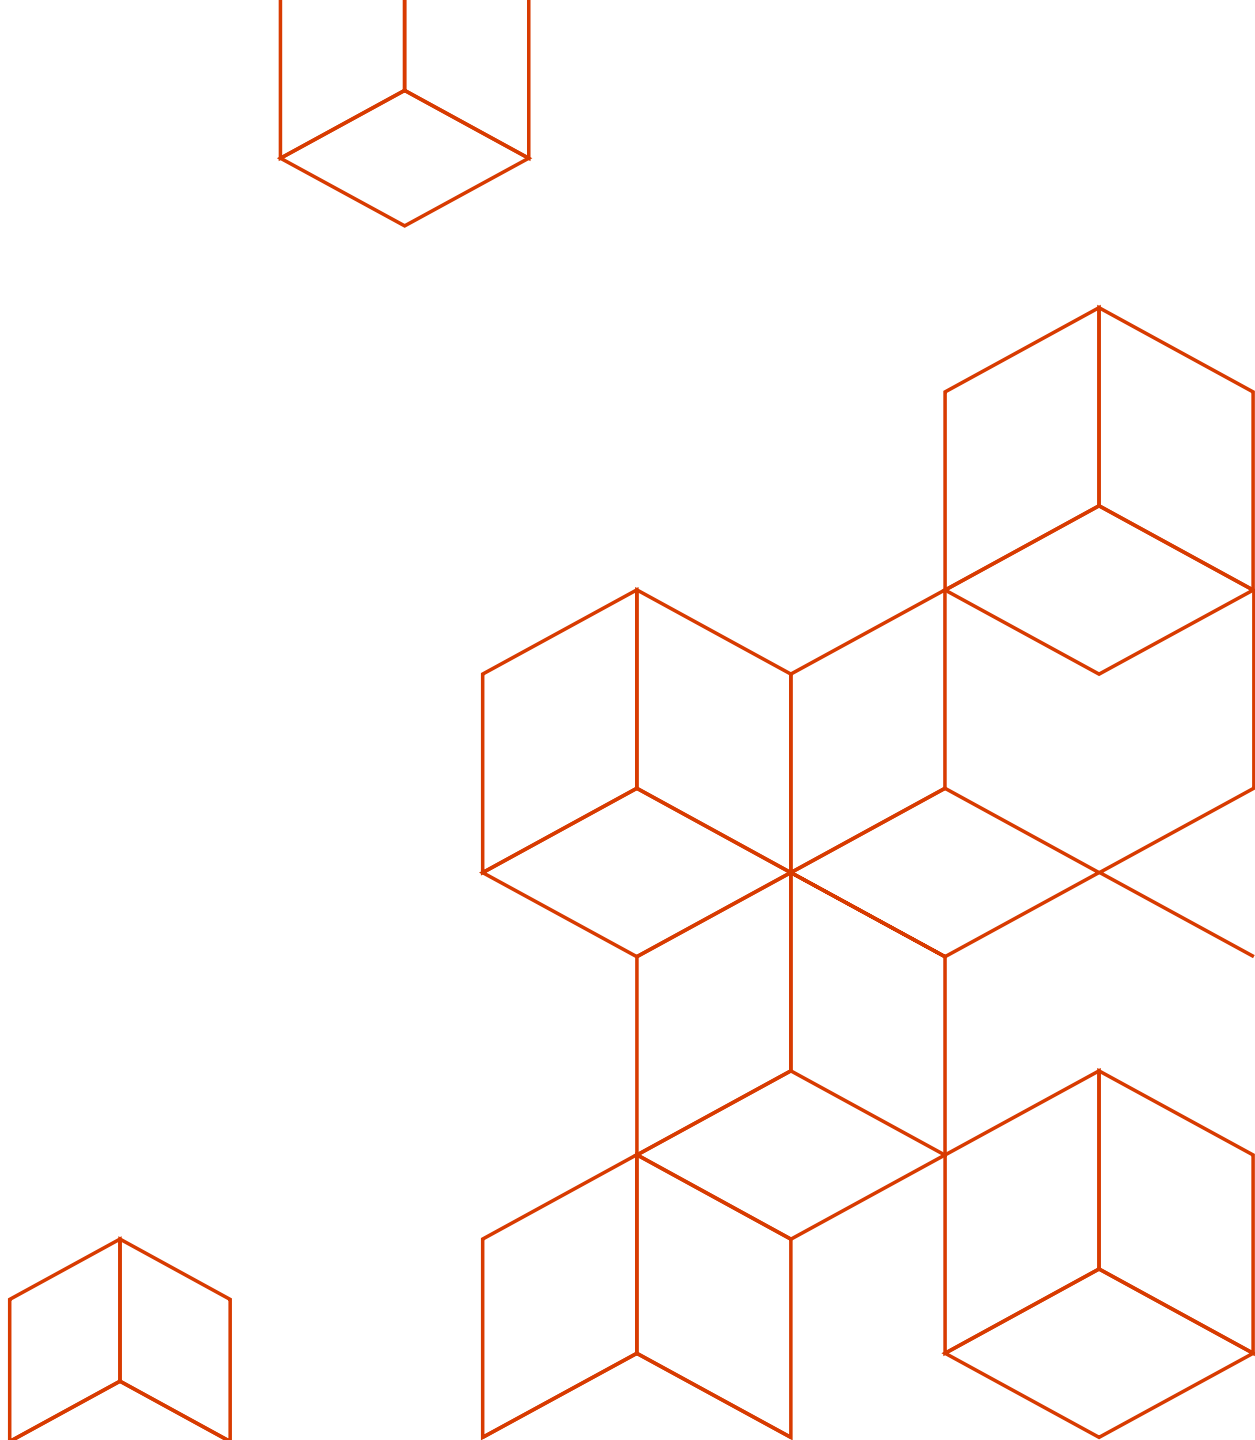

Multi-Region Clusters

- Minimize downtime risk
- One live region
  - Another backup
  - Or weighted traffic
- A/B testing

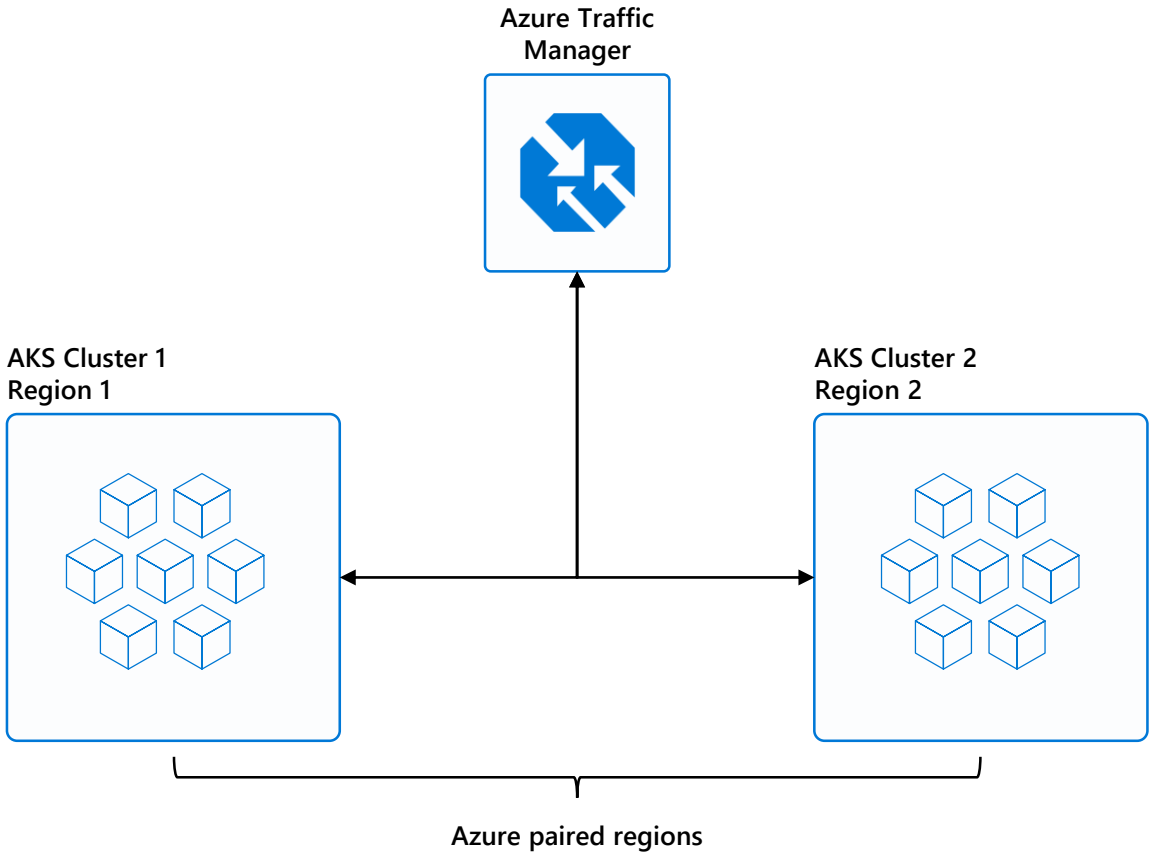

Logging and Monitoring

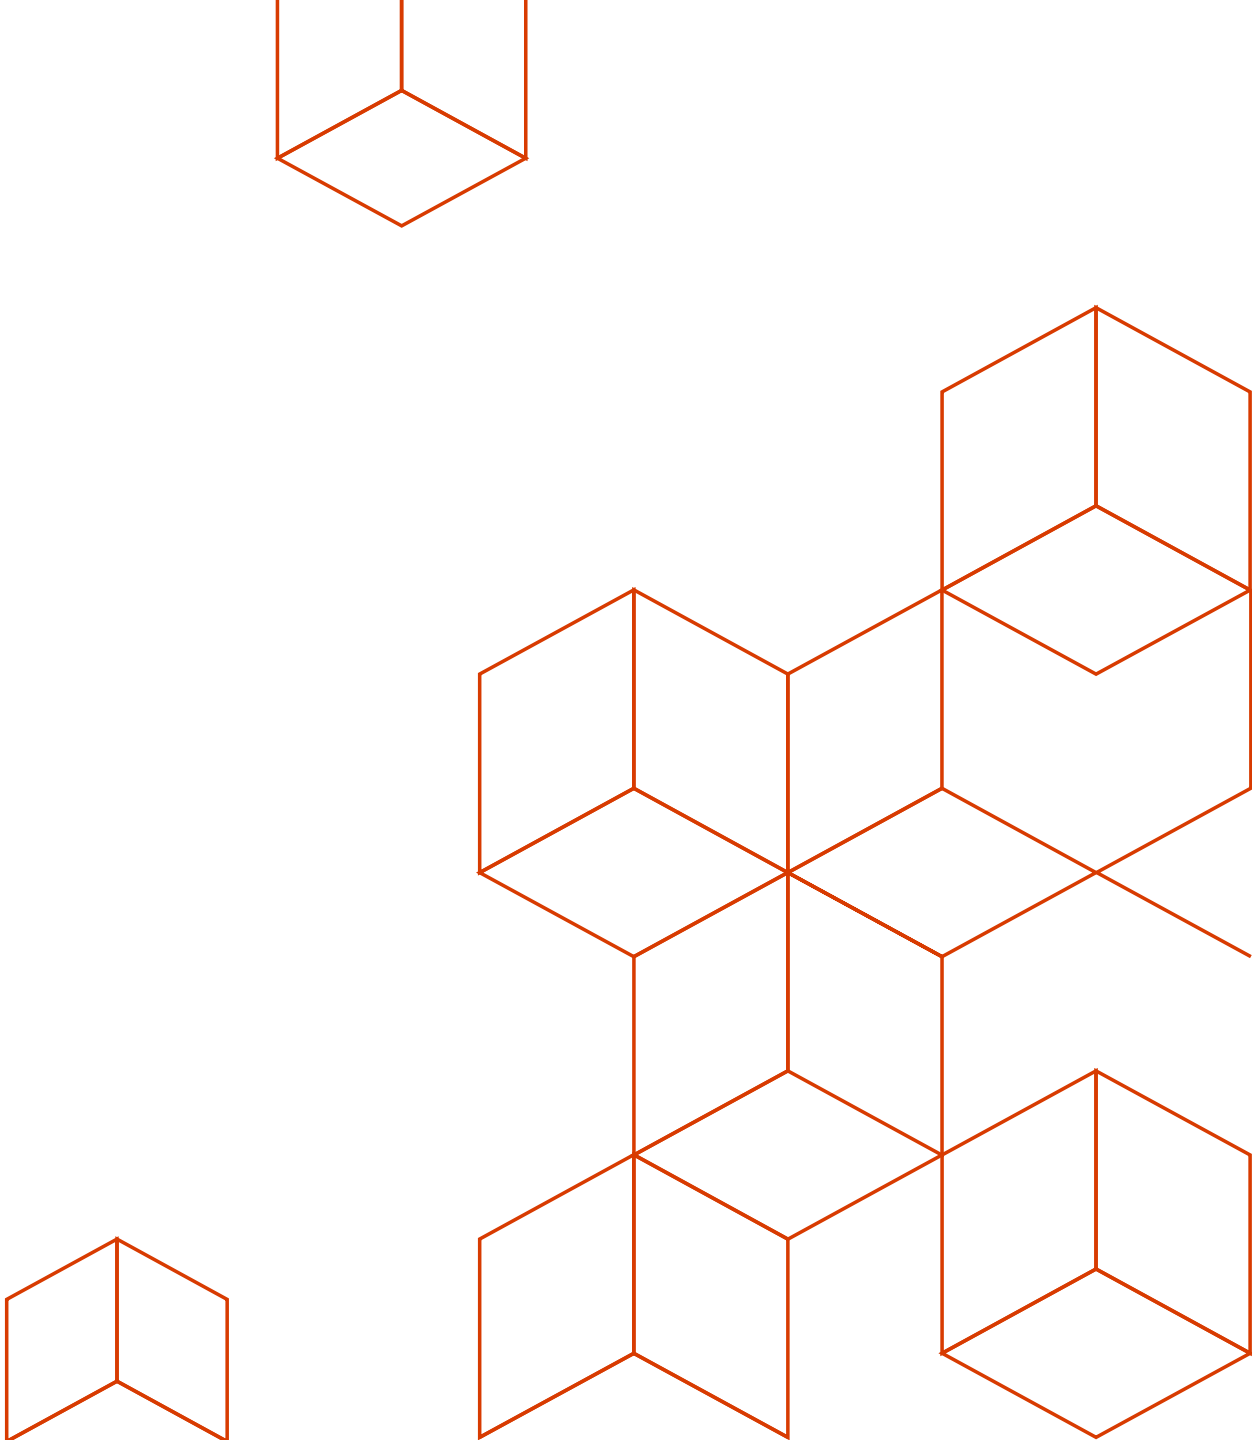

Monitoring/Logging your cluster

- Log Everything to stdout / stderr
- Key Metrics:
  - Node metrics (CPU Usage, Memory Usage, Disk Usage, Network Usage)
  - Kube_node_status_condition
  - Pod memory usage / limit; memory_failures_total
    - container_memory_working_set_bytes
  - Pod CPU usage average / limit
  - Filesystem Usage / limit
  - Network receive / transmit errors
- Azure Monitor for Containers

In the roadmap

Overview health of AKS cluster

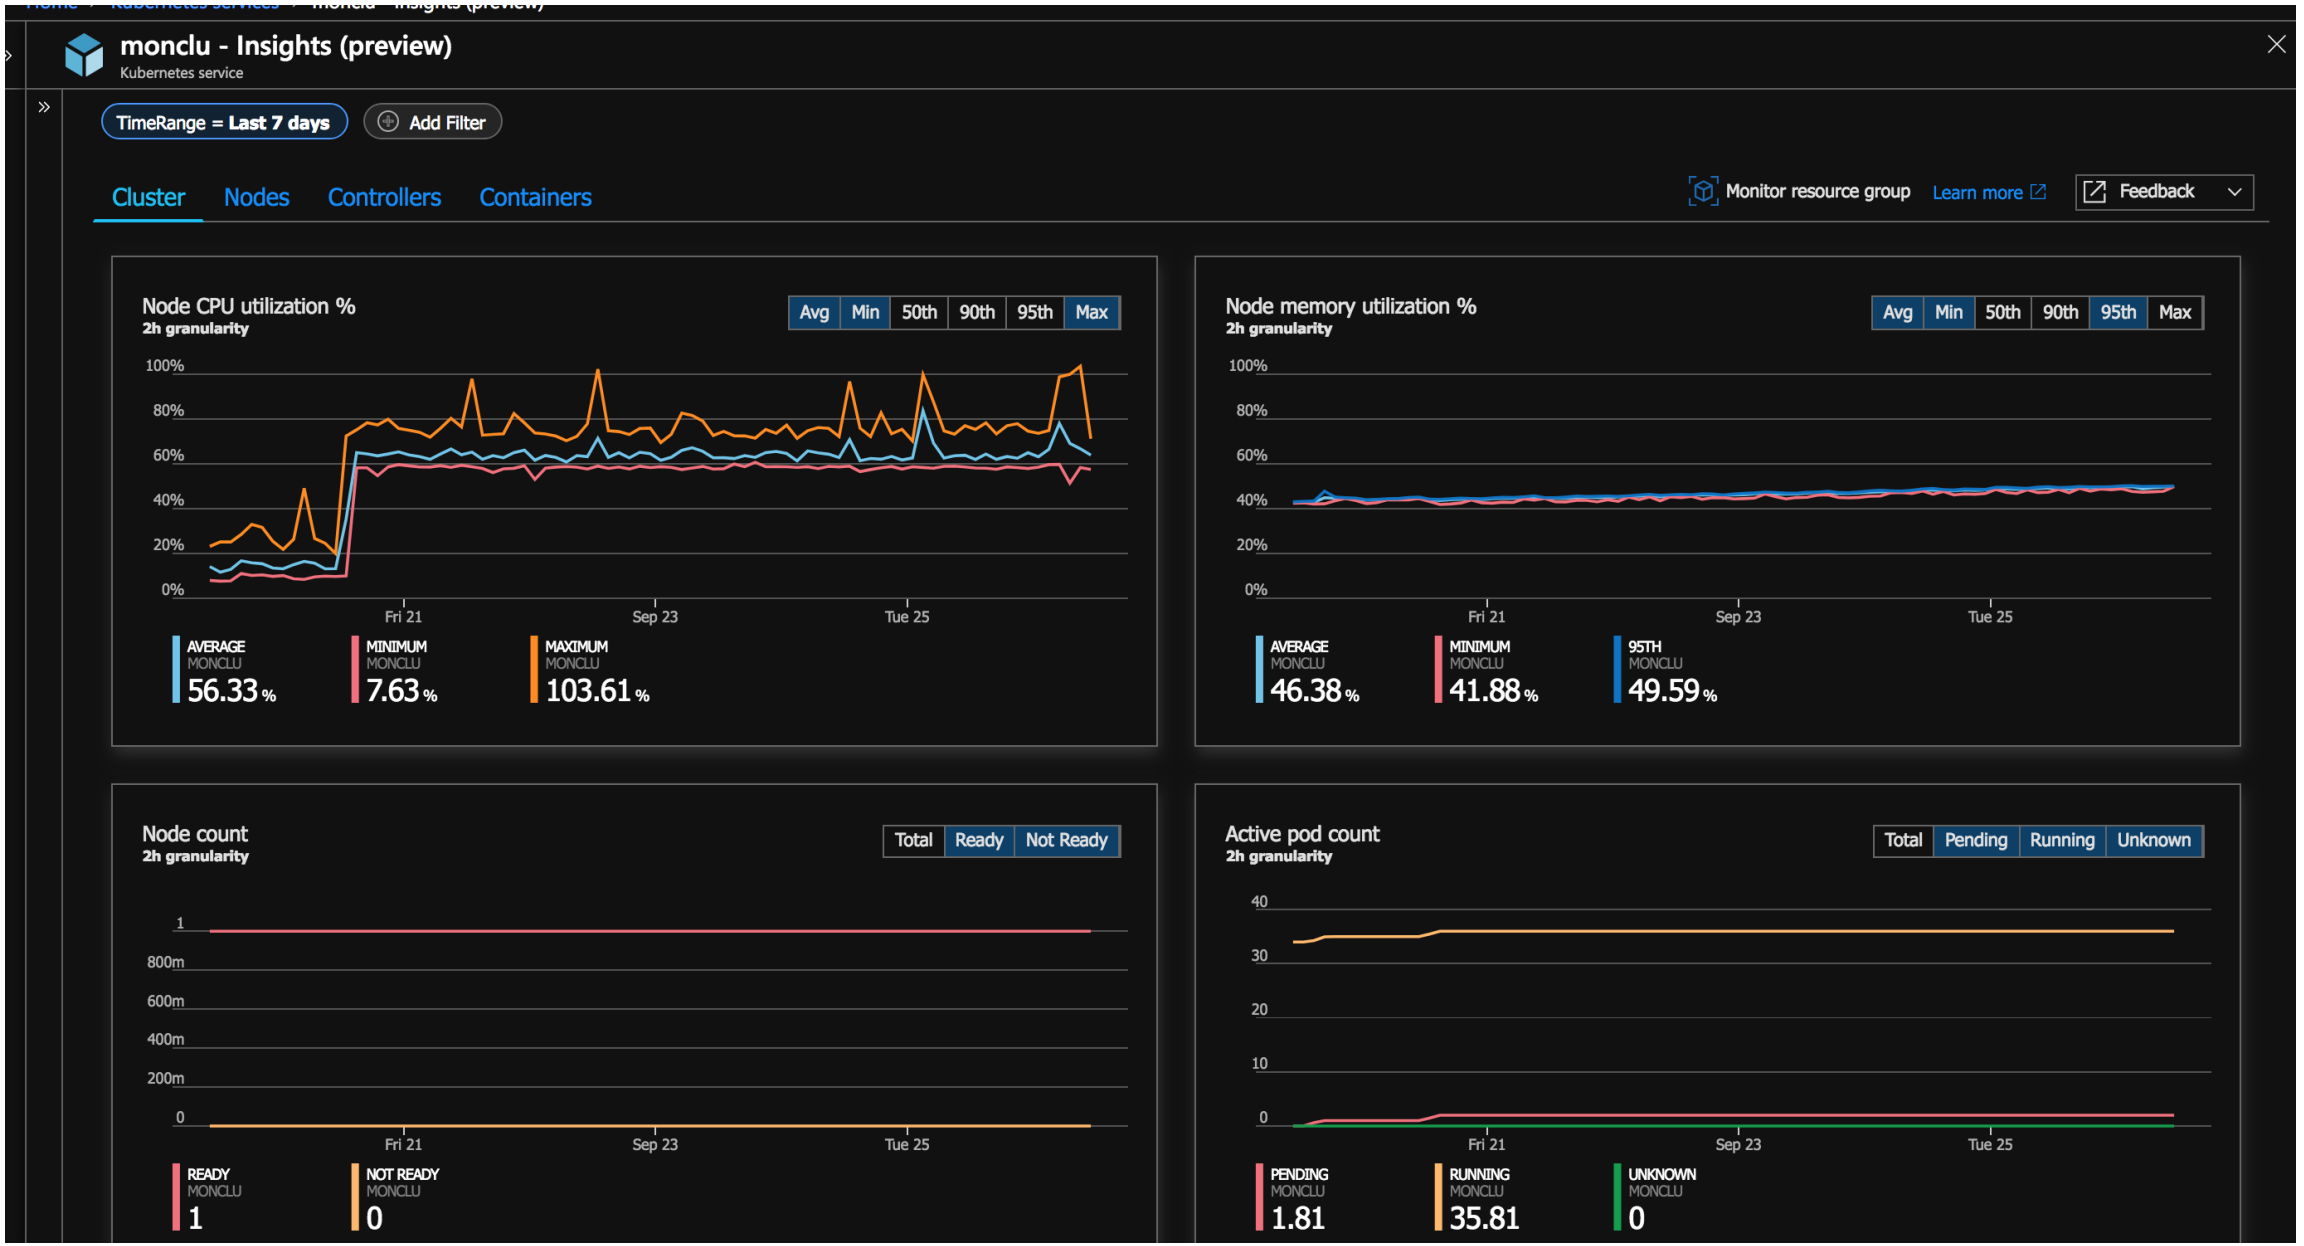

Node event Logs

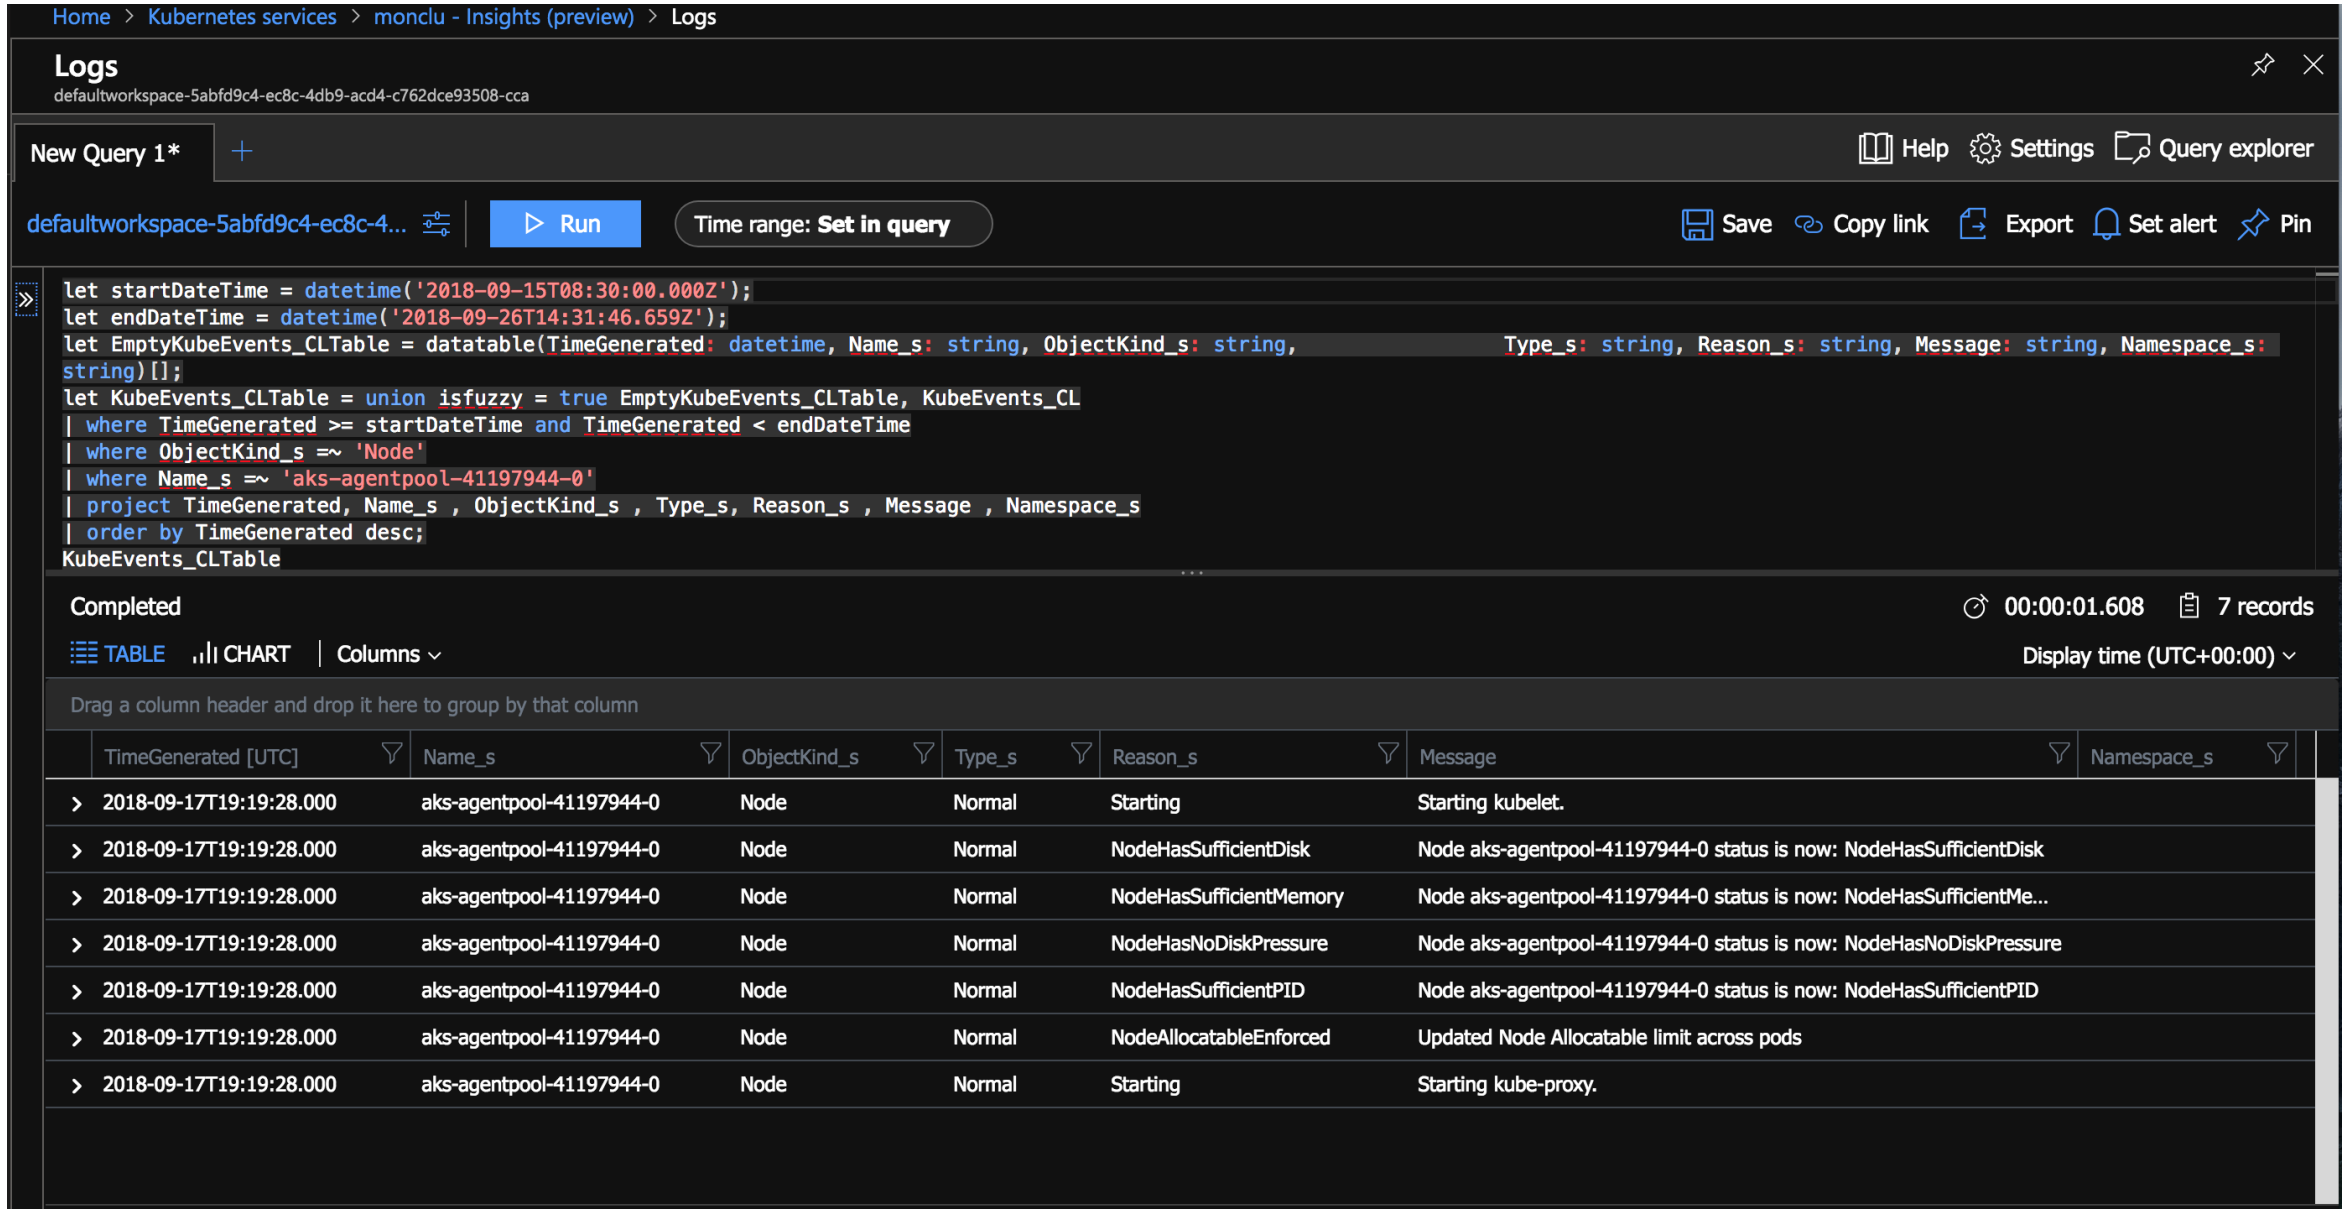

Pod usage and details

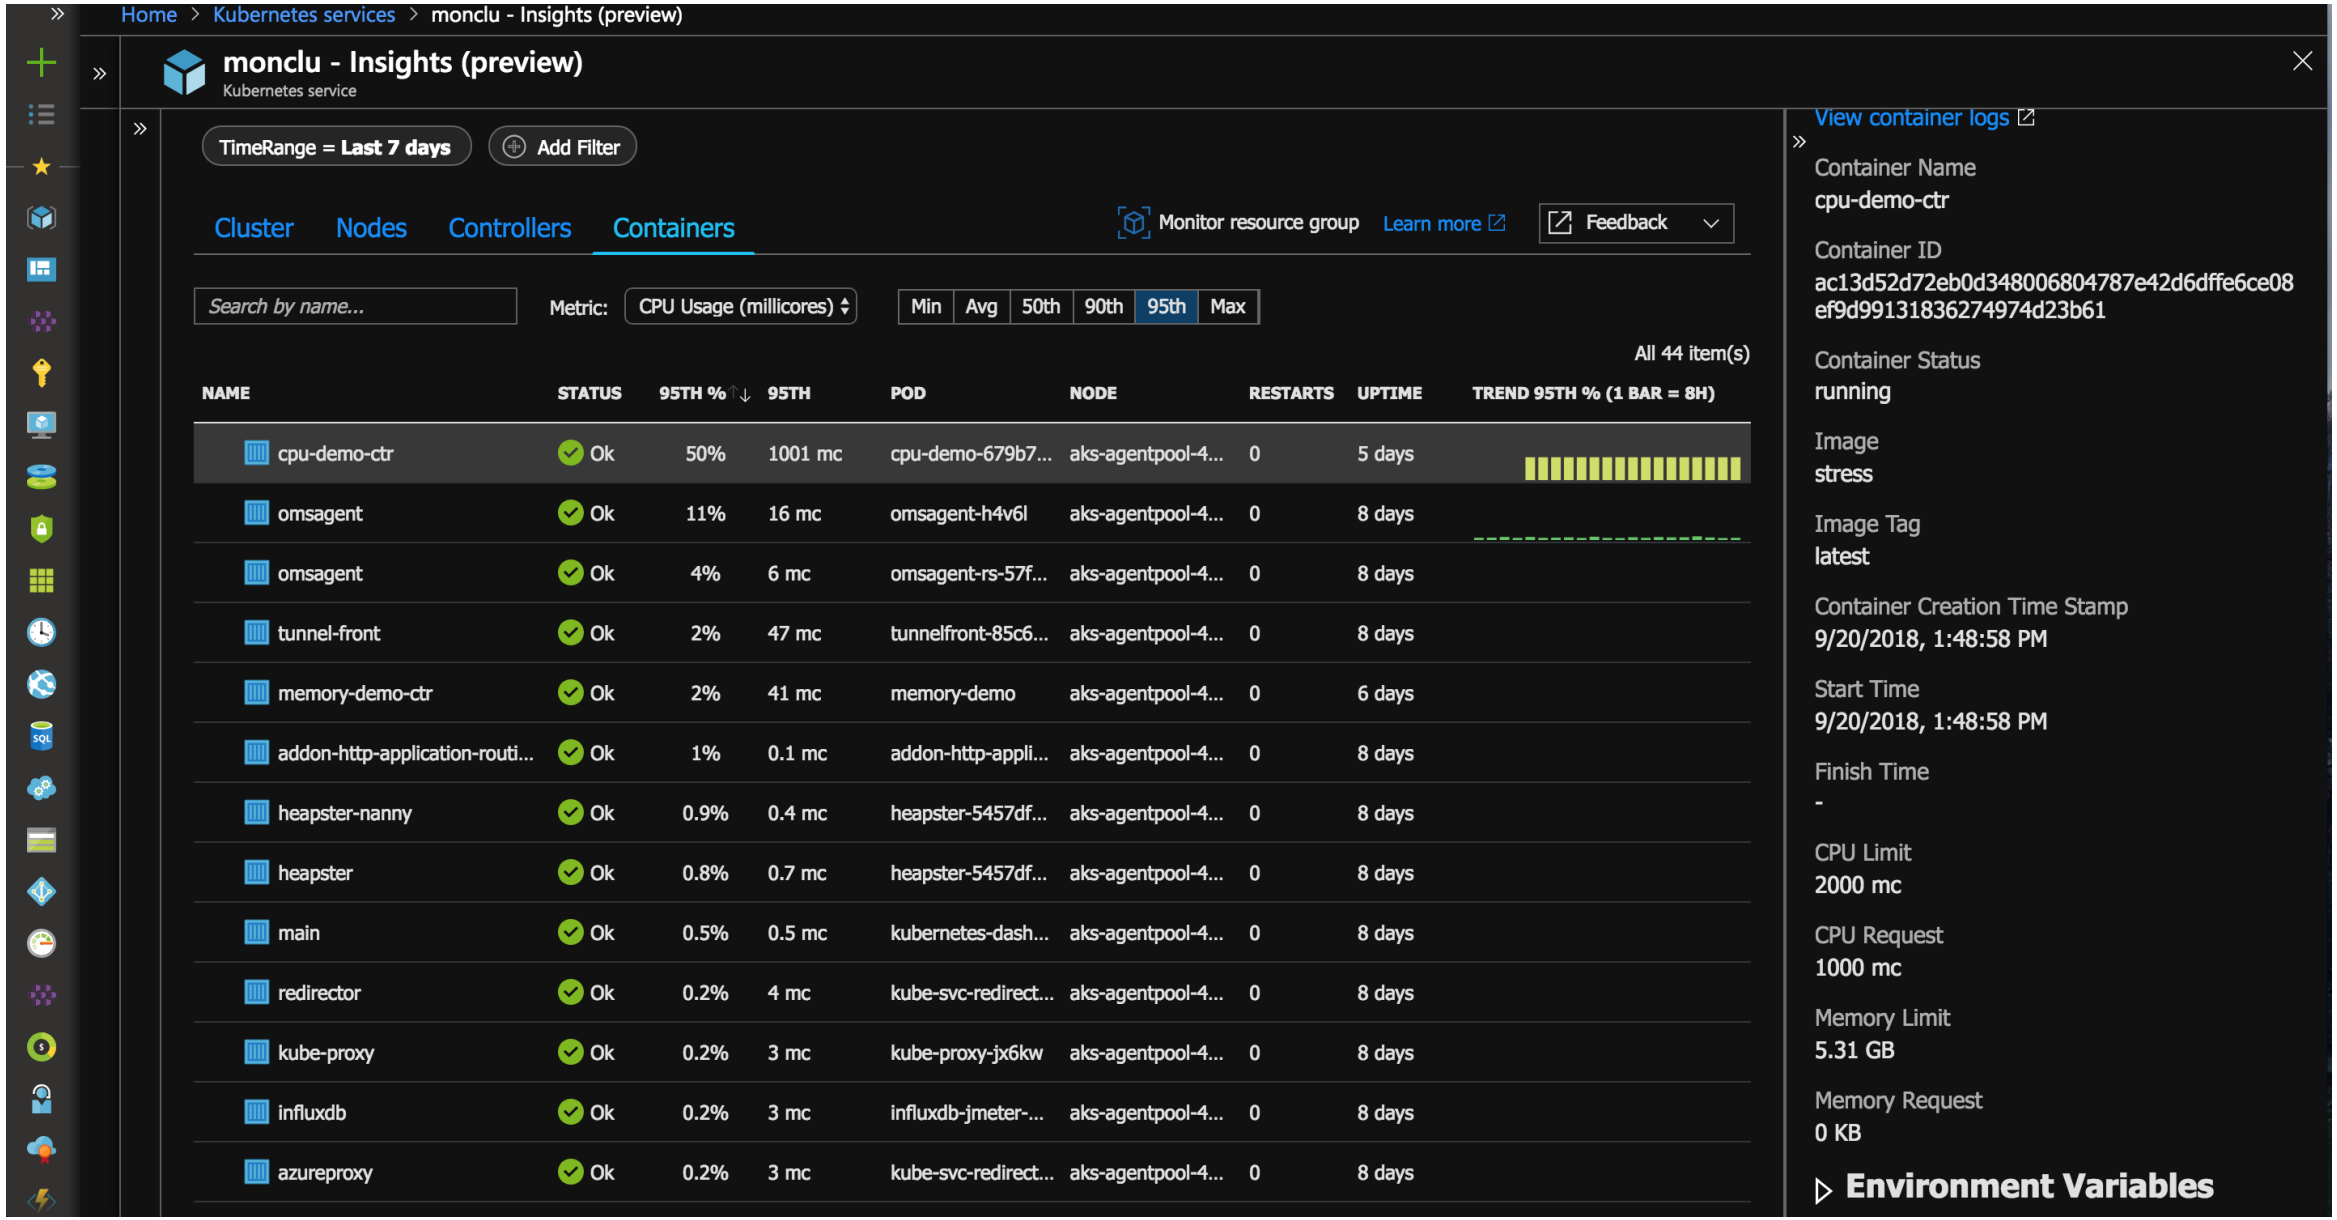

Customer control plane logs

- Use the Azure portal to enable diagnostics logs
- Pipe logs to log analytics, event hub or a storage account
- Metrics available today
  - Kube-controller-manager
  - Kube-api-server
  - Kube-scheduler
  - Audit logs on the roadmap

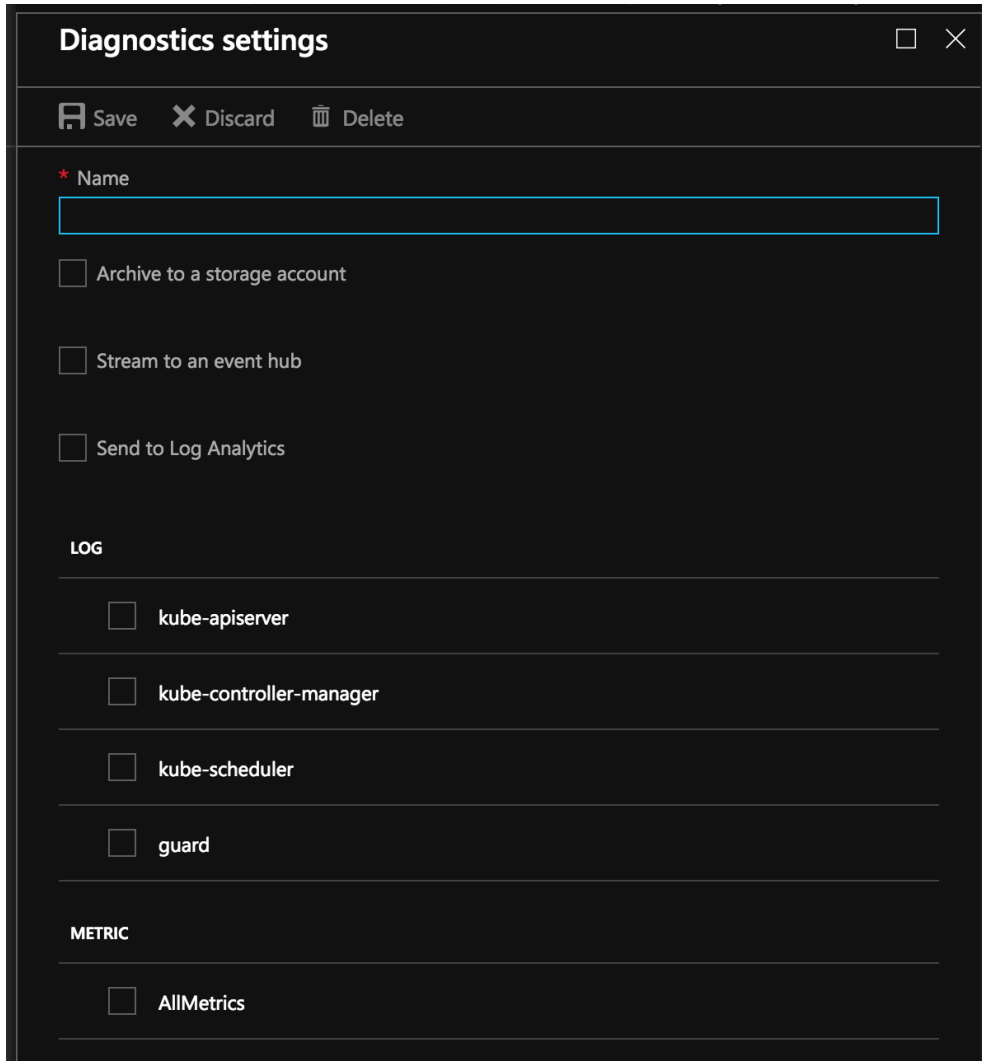

Example control plane logs

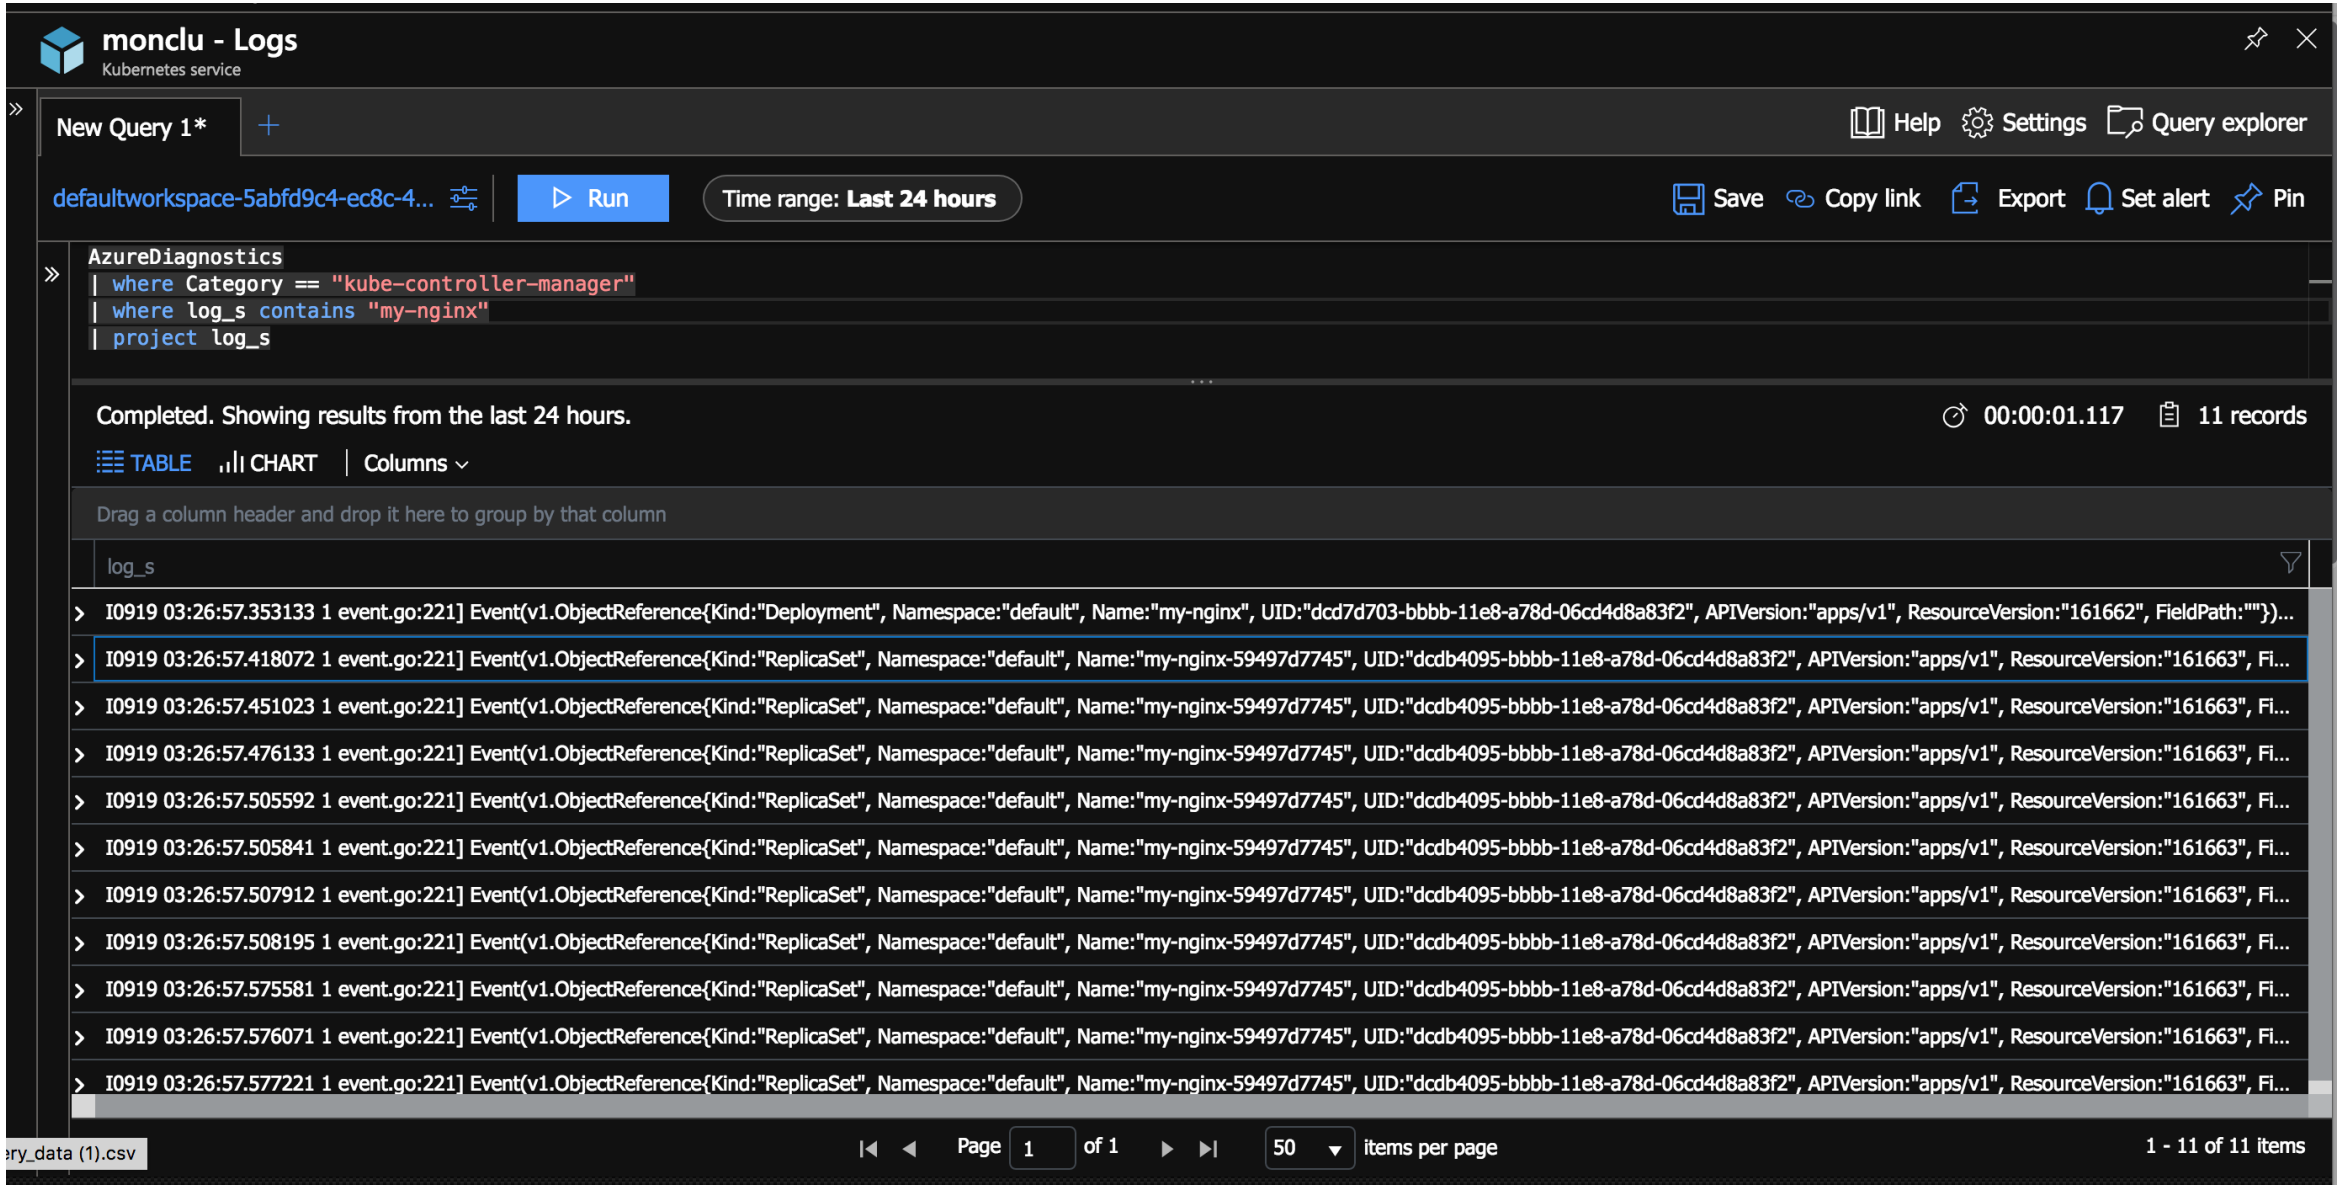

Multi cluster monitoring

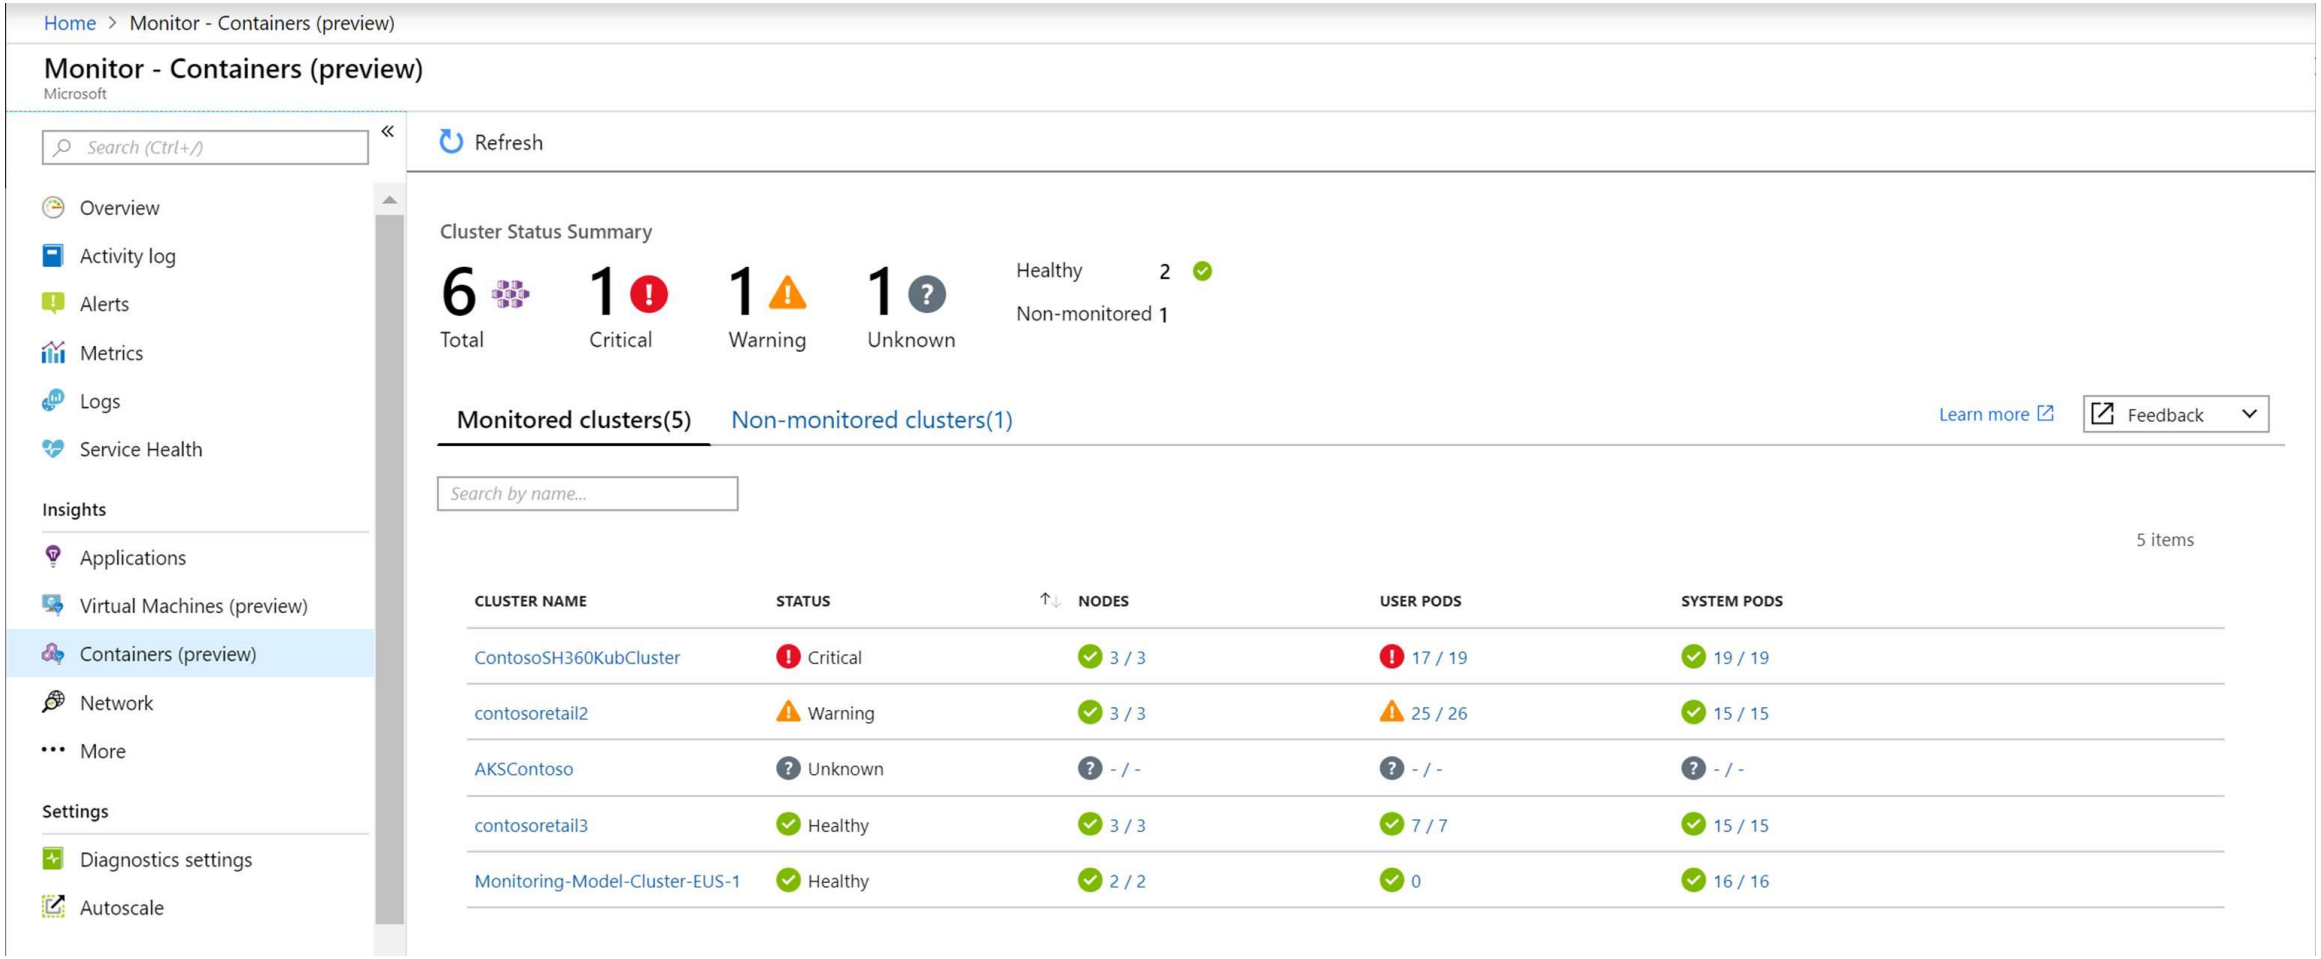

Demo

Monitoring and logging (Saurya)

1) Node/pod usage, kube events
2) Pod hogging resource, show resource request limits
3) Talk about percentile for capacity planning
4) Show filter –kube-system
5) https://aka.ms/multiaksinsights

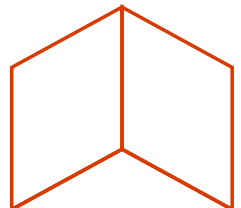

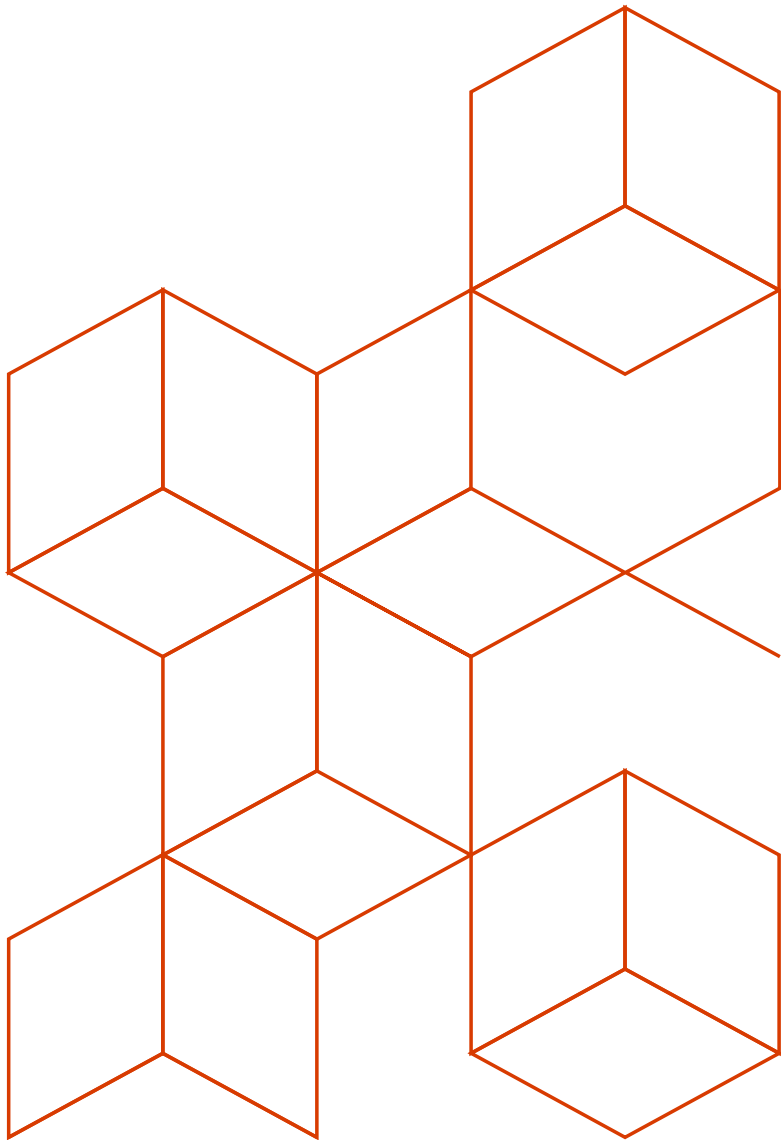

Resources

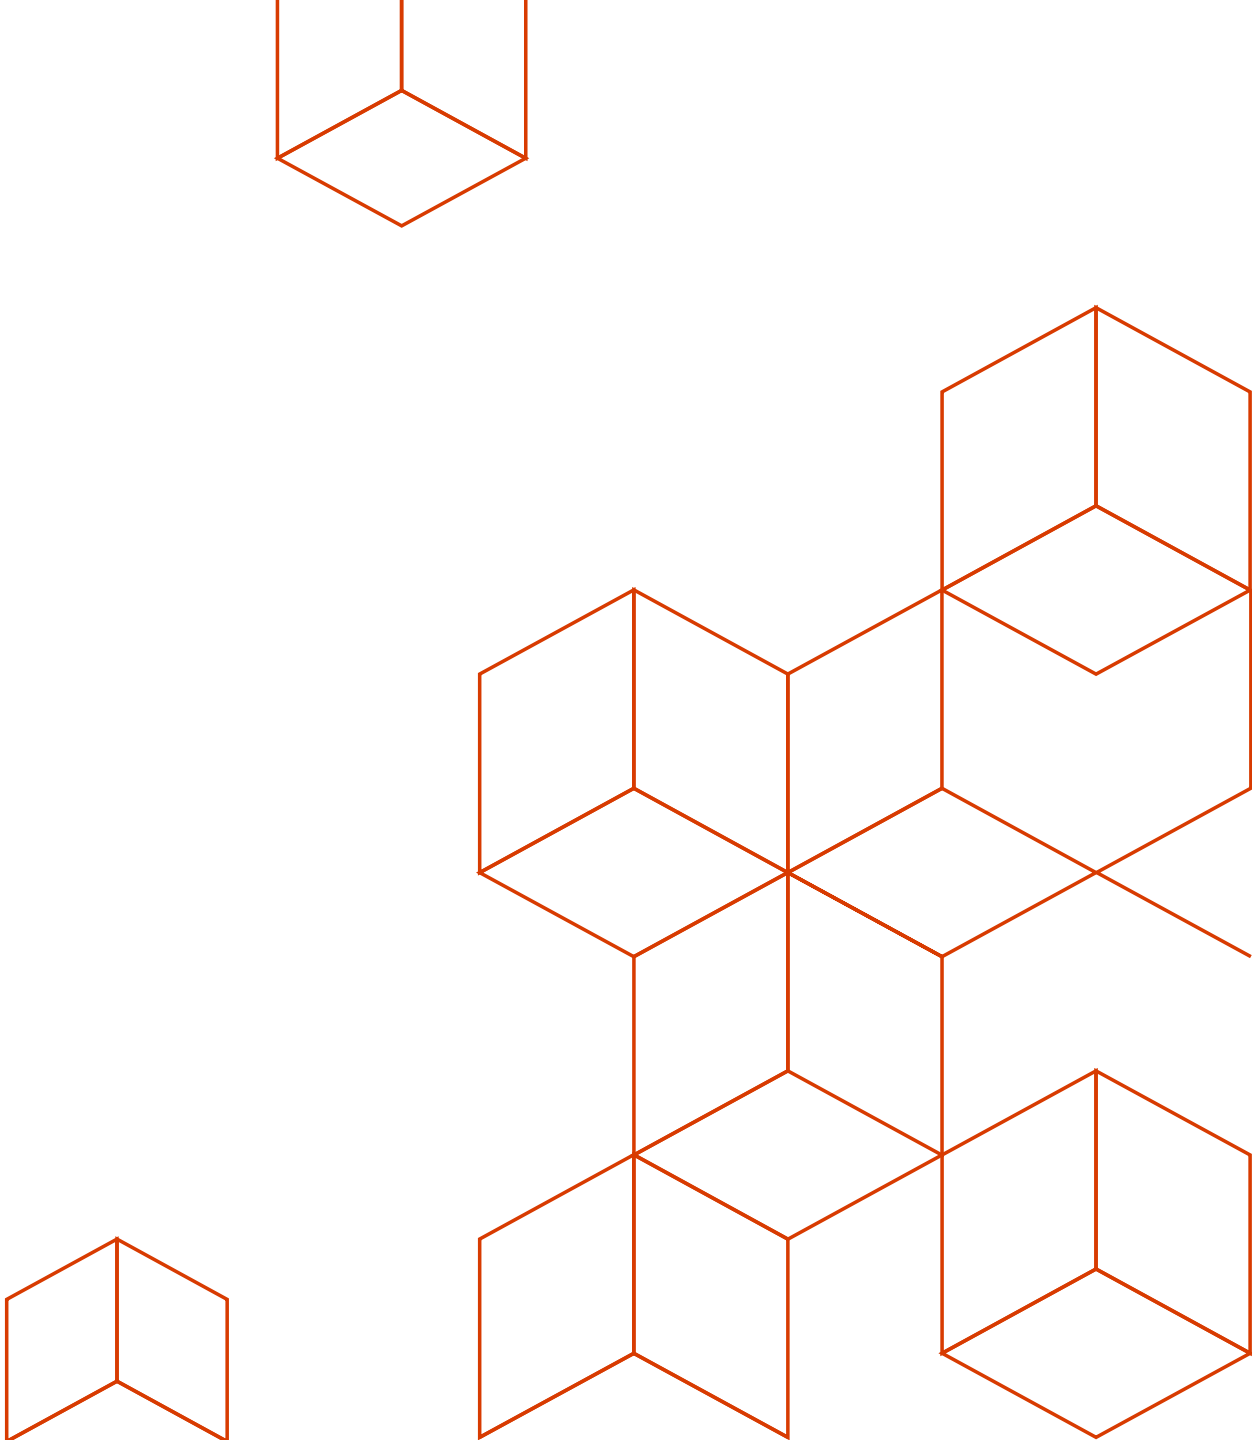

Resources

- AKS Best Practices GitHub: https://github.com/Azure/k8s-best-practices
- AKS Hackfest: aka.ms/k8s-hackfest & https://github.com/Azure/kubernetes-hackfest
- Distributed systems Labs by Brendan Burns
- Kube Advisor: https://github.com/Azure/kube-advisor
- VSCode Kubernetes Extension
- Documentation resources
- Ebook for distributed systems
- AKS HoL
- Connect with us on twitter:
  - Jorge Palma - @jorgefpalma
  - Mohammad Nofal - @mohmd_nofal
  - Saurya Das - @sauryadas
  - Andrew Randall - @andrew_randall

Thank You!

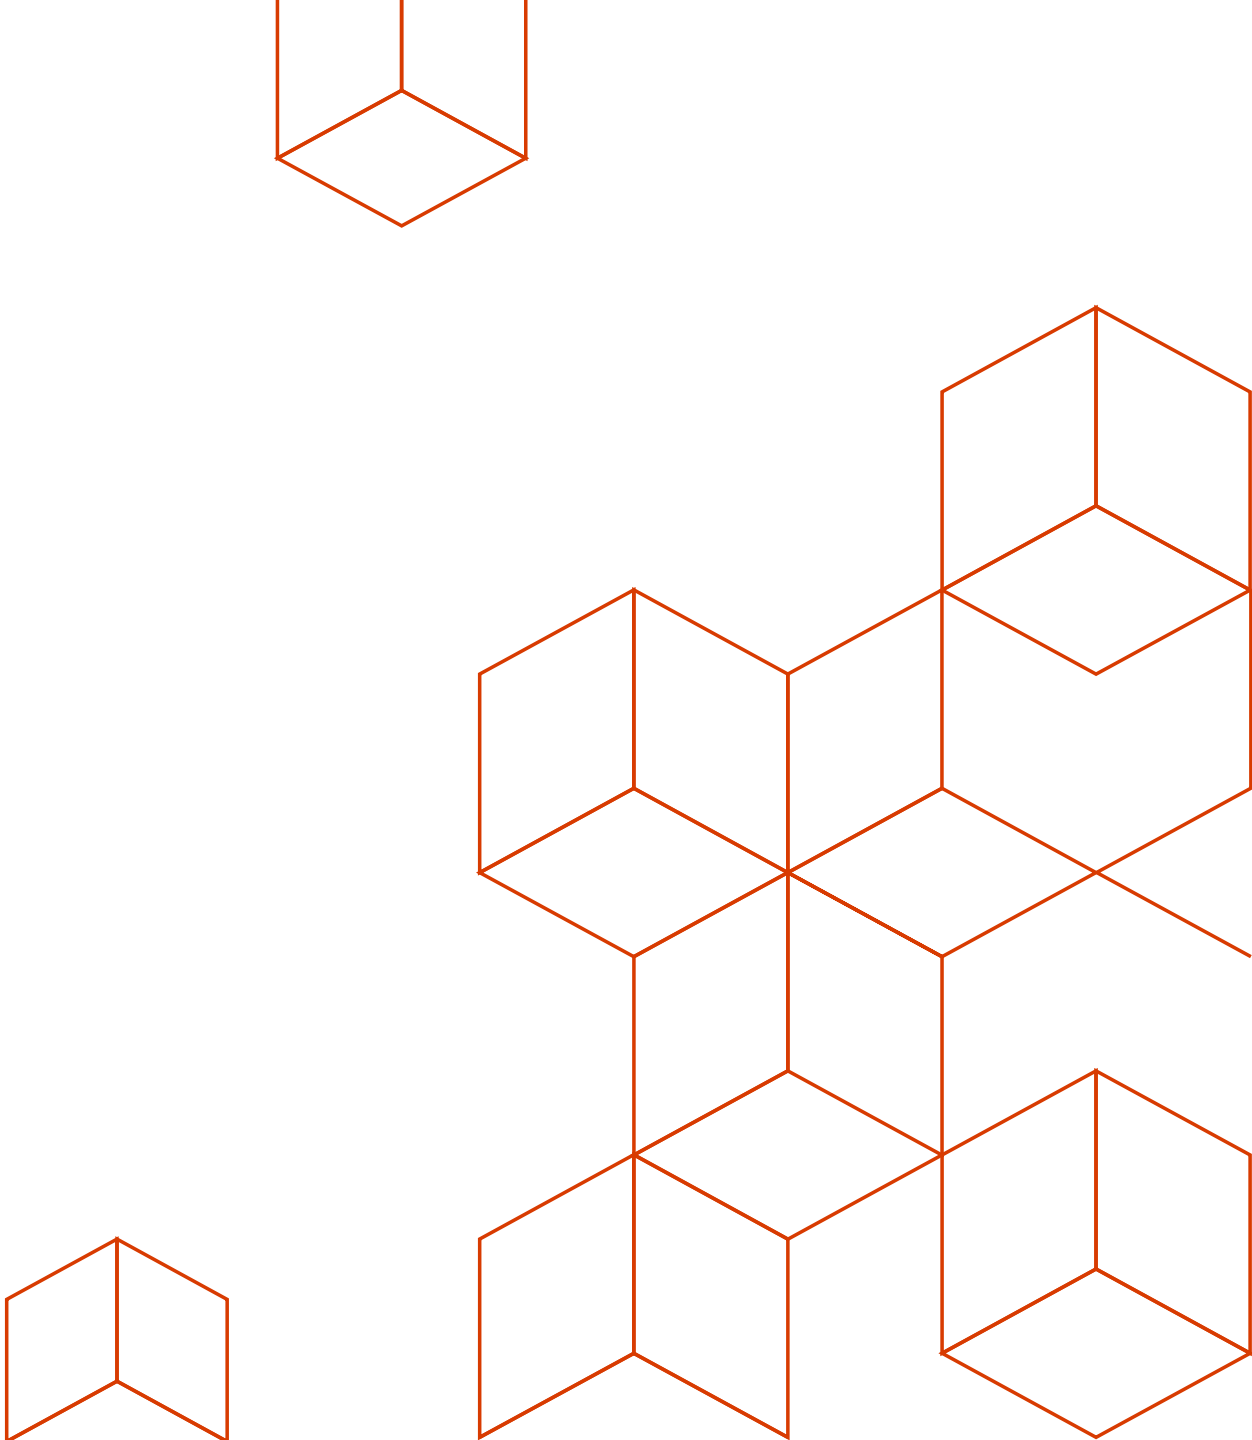

Microsoft